



Computer Science
UC Berkeley

Program Synthesis for Low-Power Accelerators

Ras Bodik
Mangpo Phitchaya Phothilimthana
Tikhon Jelvis
Rohin Shah
Nishant Totla

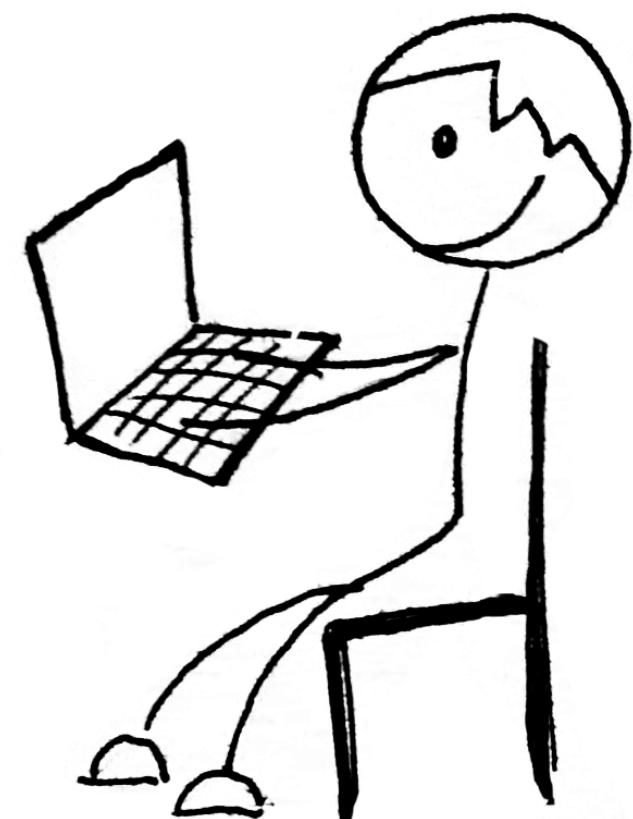
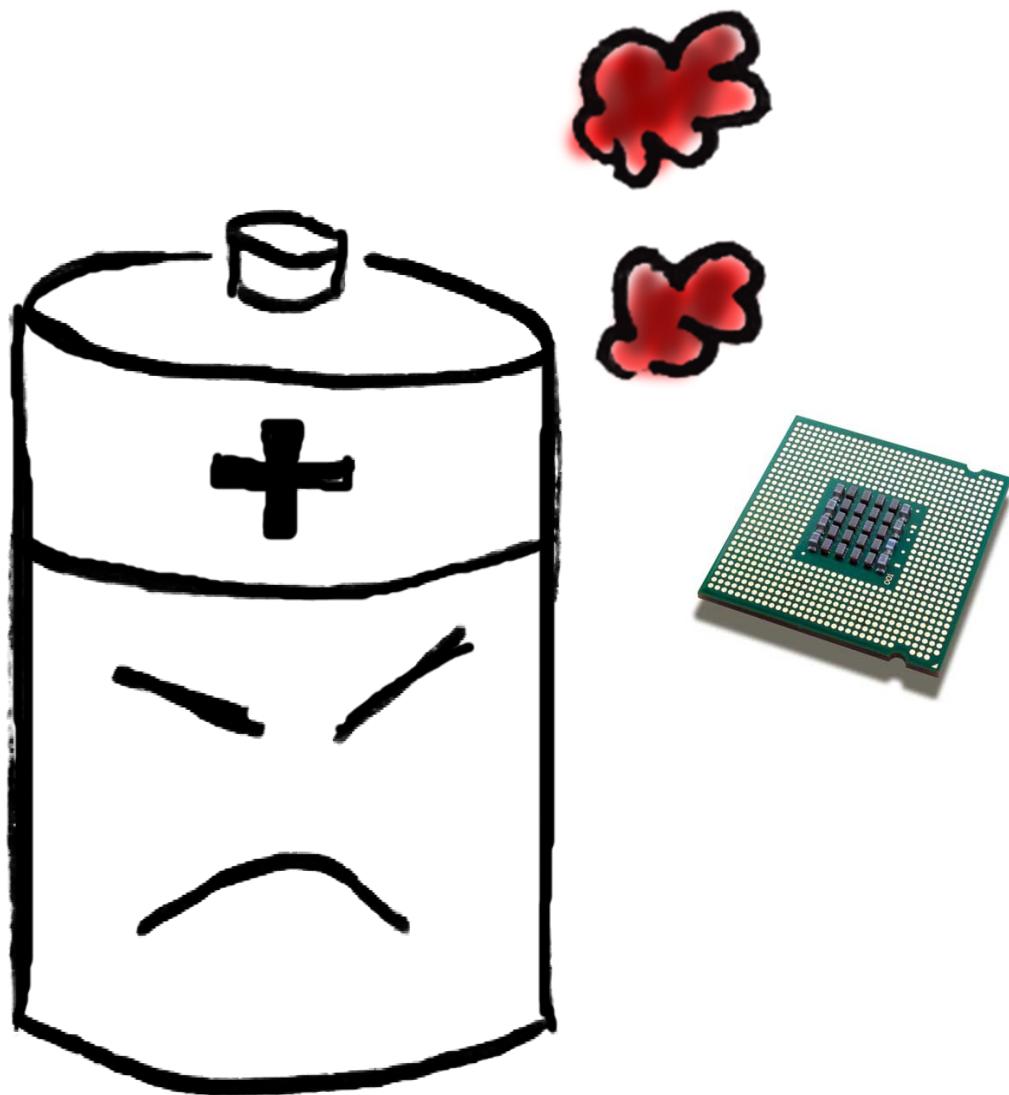
What we do and talk overview

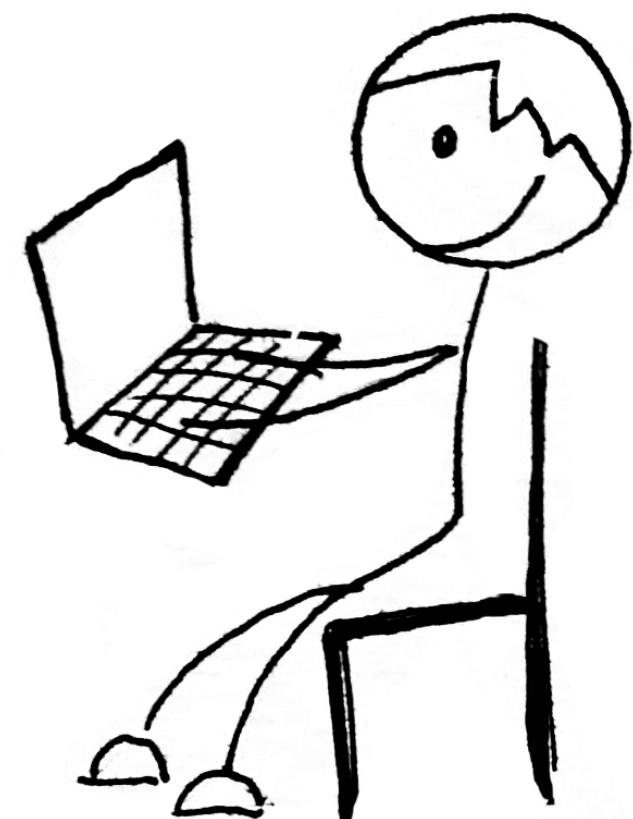
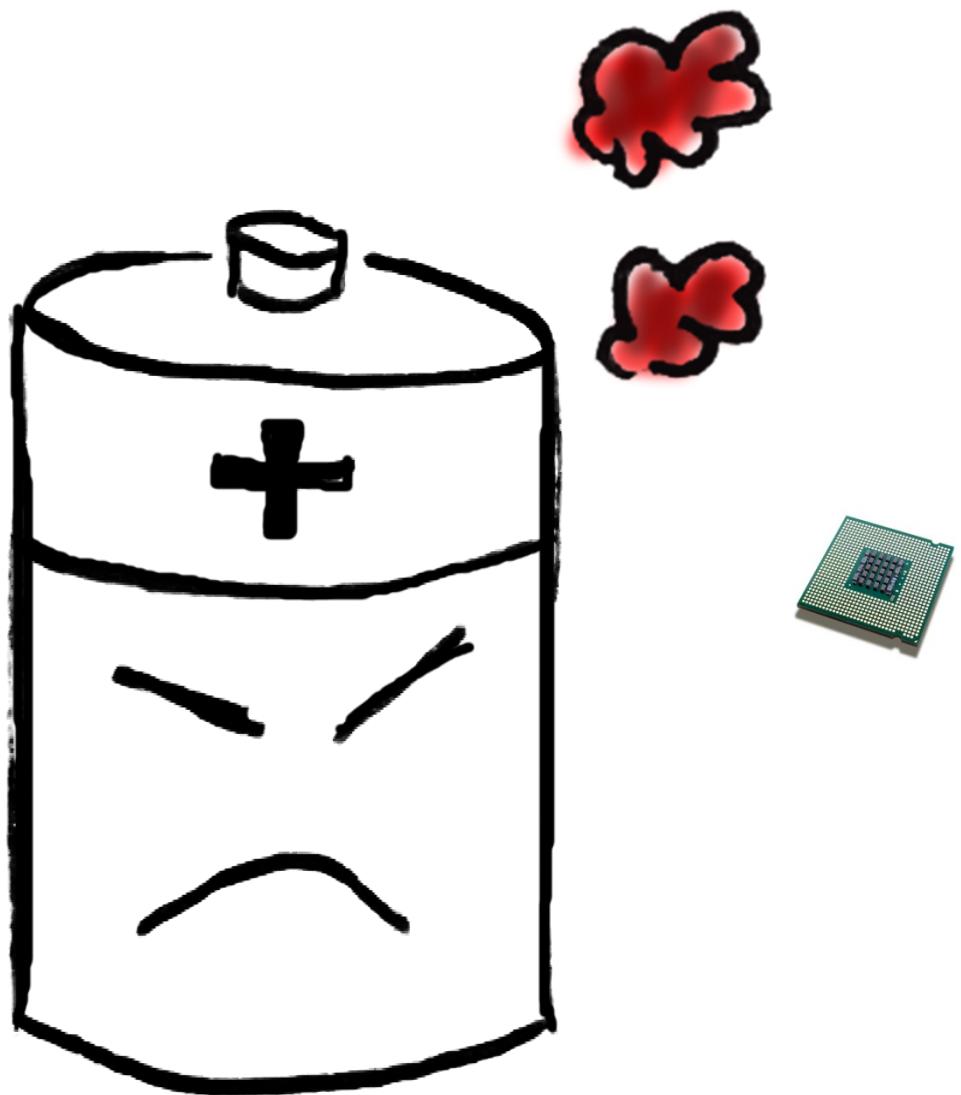
Our main expertise is in program synthesis
a modern alternative/complement to compilation

Our applications of synthesis: hard-to-write code
parallel, concurrent, dynamic programming, end-user code

In this project, we explore spatial accelerators
an accelerator programming model aided by synthesis

Future Programmable Accelerators







back to 16-bit nums

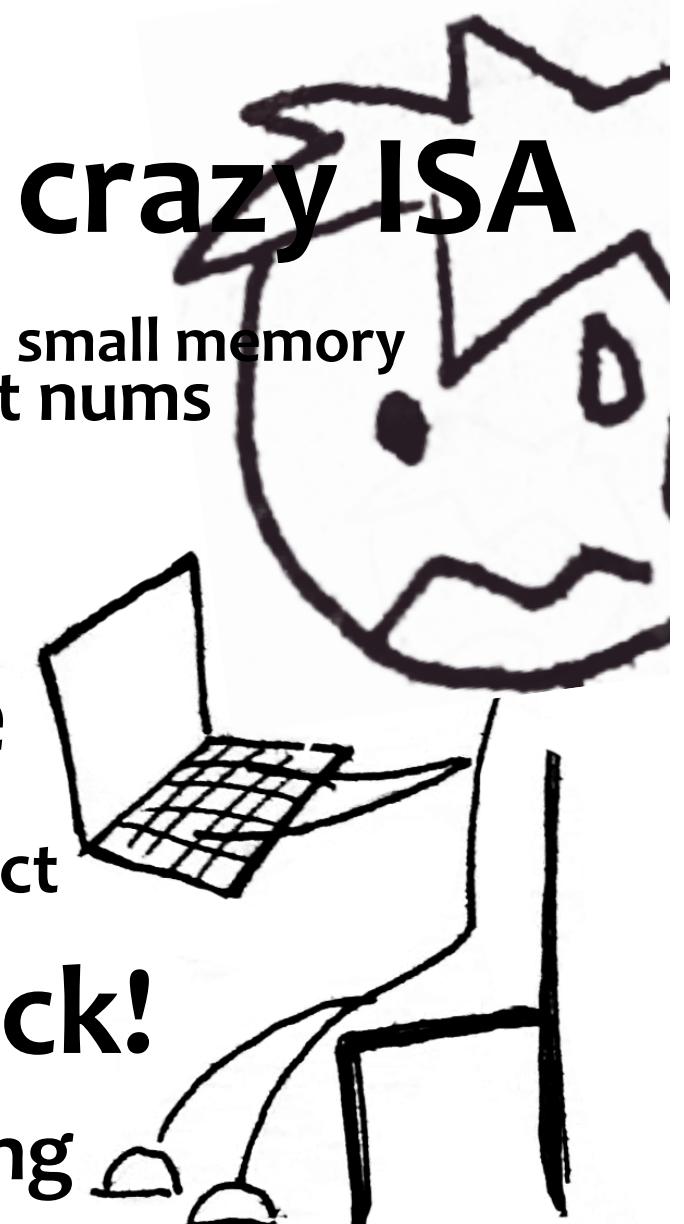


no cache coherence

limited interconnect

no clock!

spatial & temporal partitioning



crazy ISA

small memory



crazy ISA

small memory

back to 16-bit nums

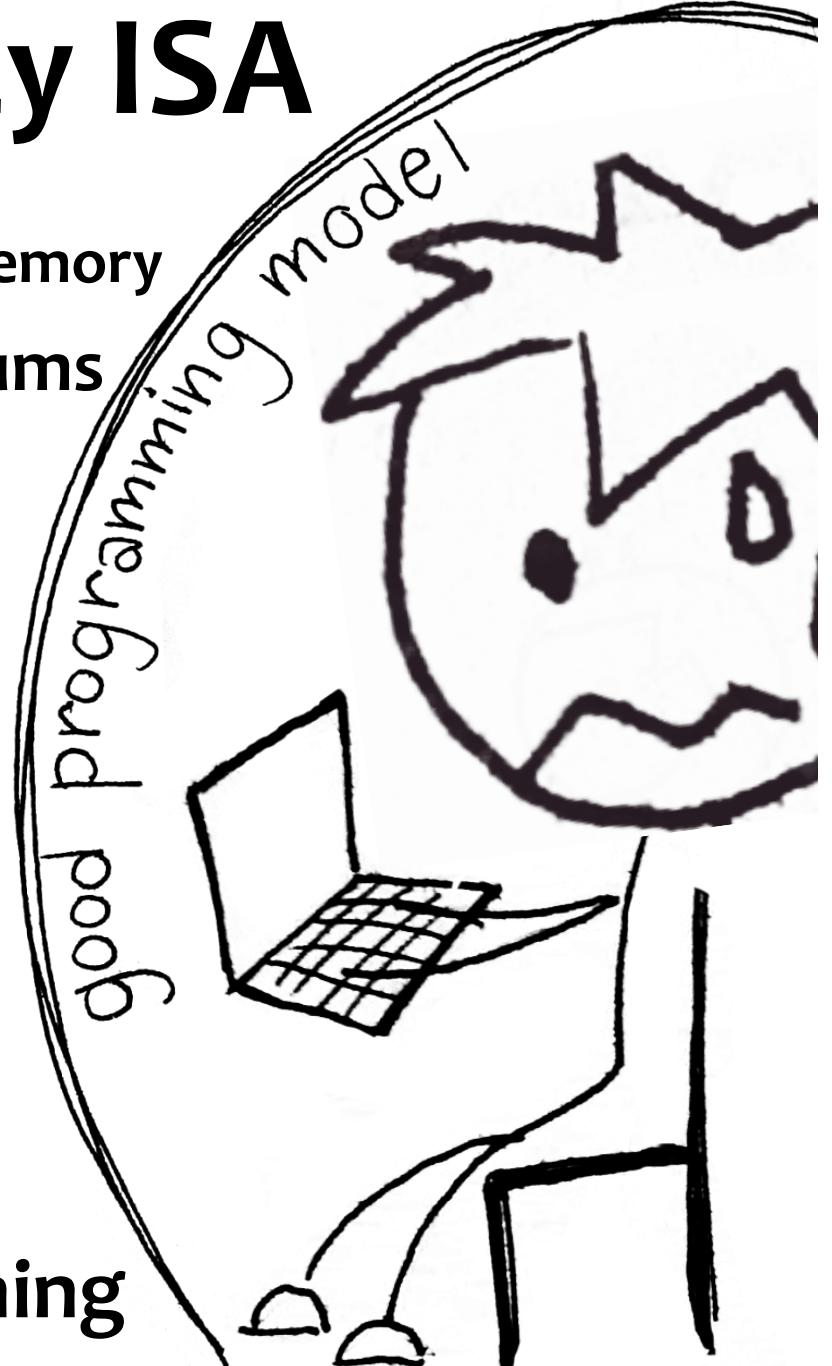


no cache coherence

limited interconnect

no clock!

spatial & temporal partitioning



What we desire from programmable accelerators

We want the obvious conflicting properties:

- high performance at low energy
- easy to program, port, and autotune

Can't usually get both

- transistors that aid programmability burn energy
- ex: cache coherence, smart interconnects, ...

In principle, most decisions can be done in compilers

- which would simplify the hardware
- but the compiler power has proven limited

We ask

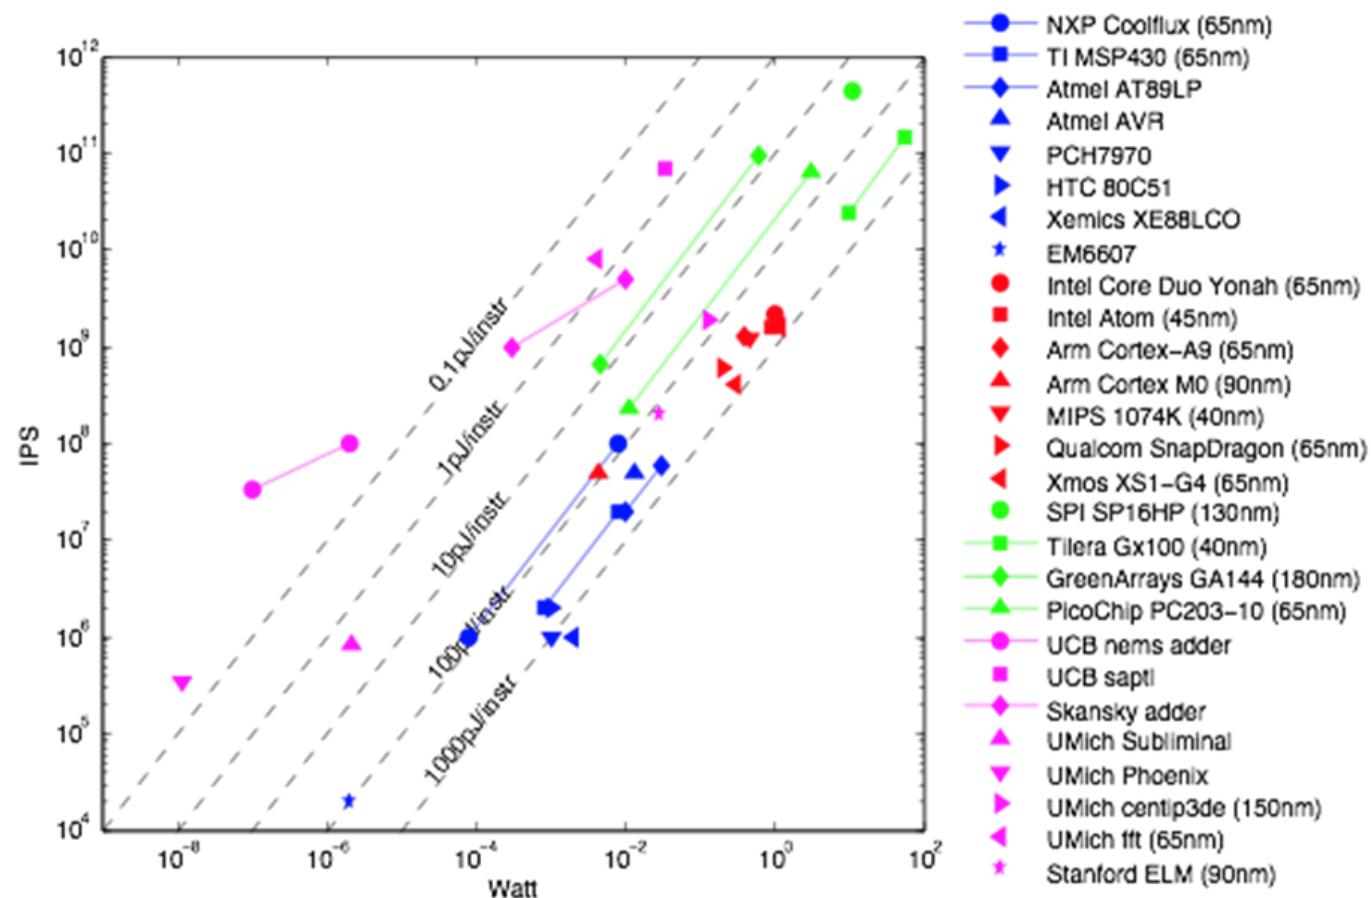
How to use fewer “programmability transistors”?

How should a programming framework support this?

Our approach: synthesis-aided programming model

GA 144: our stress-test case study

Low-Power Architectures



thanks: Per Ljung (Nokia)

Why is GA low power

The GA architecture uses several mechanisms to achieve ultra high energy efficiency, including:

- asynchronous operation (no clock tree),
- ultra compact operator encoding (4 instructions per word) to minimize fetching,
- minimal operand communication using stack architecture,
- optimized bitwidths and small local resources,
- automatic fine-grained power gating when waiting for inter-node communication, and
- node implementation minimizes communication energy.

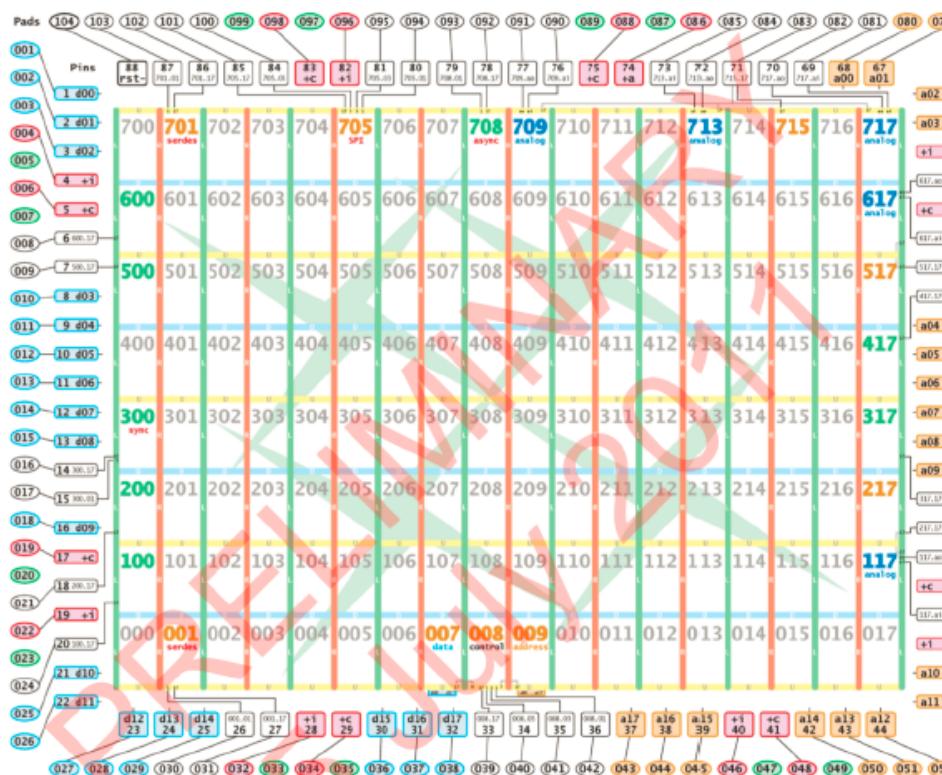
GreenArray GA144

- The GA144 is an 8 x 18 array of F18A asynchronous computers (cores)
- Cores are independent – no shared resources (i.e. clock, cache, memory bus)
 - Only globally routed signals are power/ground and reset! Very scalable architecture.
- Each core can (only!) communicate with its neighbors to the NSEW
- Cores located along perimeter include I/O functionality

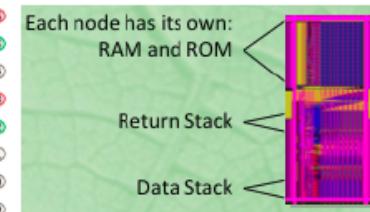
Implemented
in 180nm process
VDD = 1.8V

Instantaneous
power usage
ranges from
14 uW - 650 mW

Adjacent cores
are mirror images
of each other



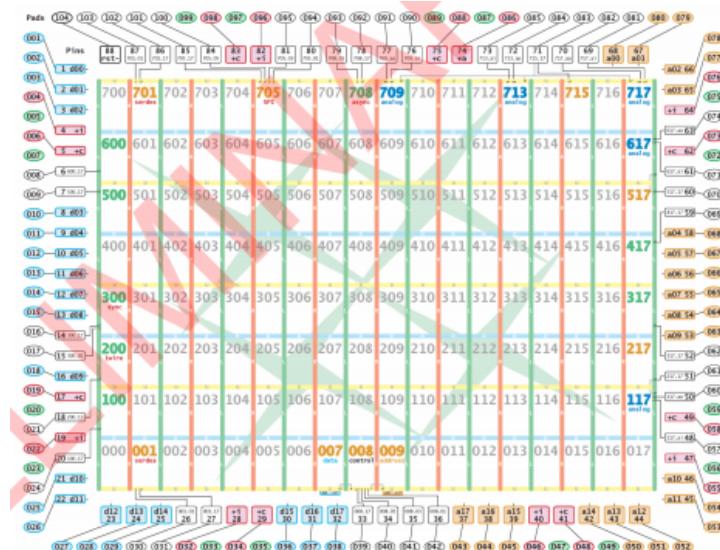
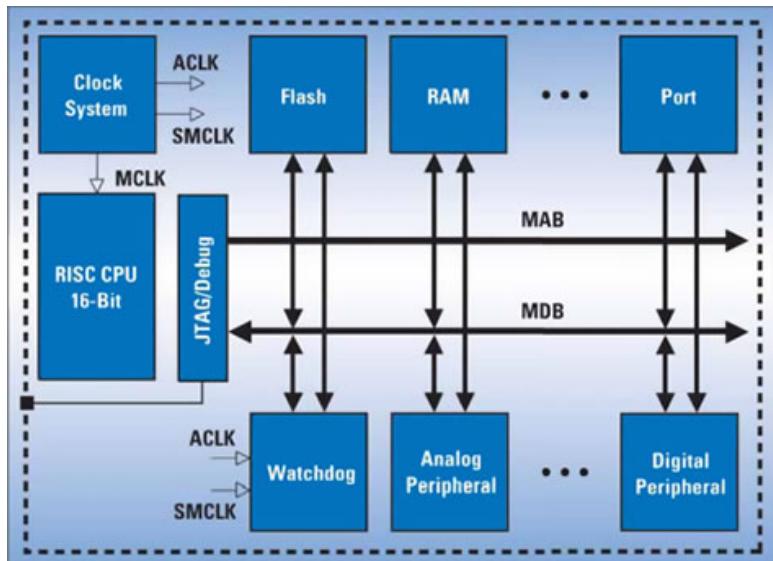
F18A layout



Fewer than 20k
transistors per
core!

Total on-chip RAM =
9216 18-bit words
(equivalent amount
of ROM)

MSP430 vs GreenArray



Finite Impulse Response Benchmark

Performance	MSP430 (65nm)			GA144 (180nm)
	F2274 swmult	F2617 hwmult	F2617 hwmult+dma	
usec / FIR output	688.75	37.125	24.25	2.18
nJ / FIR output	2824.54	233.92	152.80	17.66

MSP430 **GreenArrays**

GreenArrays 144 is **11x faster** and simultaneously
9x more energy efficient than MSP 430.

How to compile to spatial architectures

Programming/compiling challenges

Partition the code and map it to cores

also, design the algorithm for optimal partitioning

Implement inter-core communication

don't introduce races and deadlocks

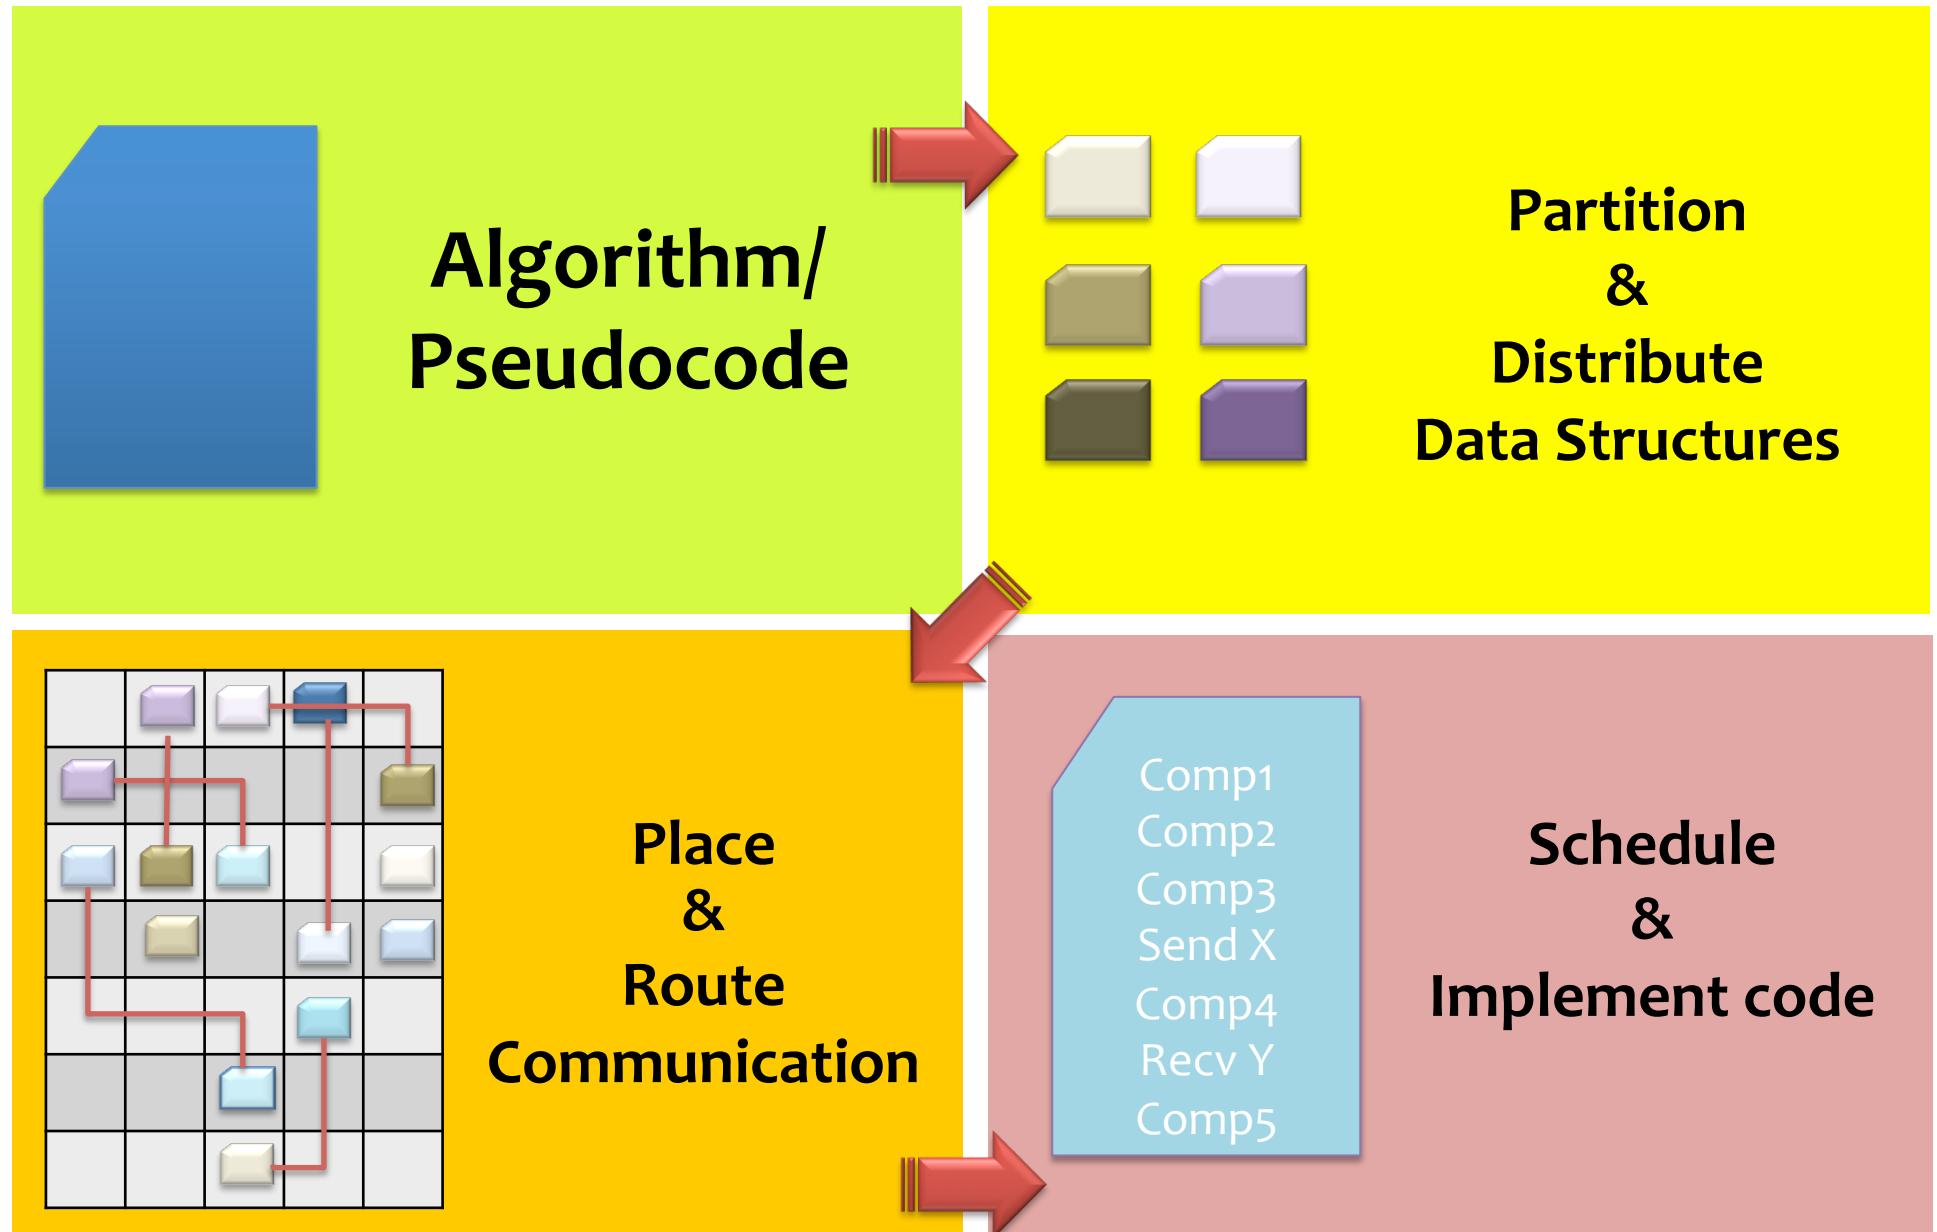
Manage distributed data structures

for user-facing accelerators (GUIs), it's not just arrays

Generate efficient code

often on a limited ISA

Programming/Compiling Challenges



Existing technologies

Optimizing compilers

hard to write optimizing code transformations

FPGAs: synthesis from C

partition C code, map it and route the communication

Bluespec: synthesis from C or SystemVerilog

synthesizes HW, SW, and HW/SW interfaces

What does our synthesis differ?

FPGA, BlueSpec:

partition and map to minimize cost

We push synthesis further:

- super-optimization: synthesize code that meets a spec
- a different target: accelerator, not FPGA, not RTL

Benefits:

superoptimization: easy to port as there is no need to port code generator and the optimizer to a new ISA

Overview of program synthesis

Synthesis with “sketches”

Extend your language with two constructs

spec:

```
int foo (int x) {  
    return x + x;  
}
```

$$\phi(x,y):y=foo(x)$$

sketch:

```
int bar (int x) implements foo {  
    return x << ??;  
}
```

?? substituted with an
int constant meeting ϕ

result:

```
int bar (int x) implements foo {  
    return x << 1;  
}
```

instead of **implements**, assertions over safety properties can be used

Example: 4x4-matrix transpose with SIMD

a functional (executable) specification:

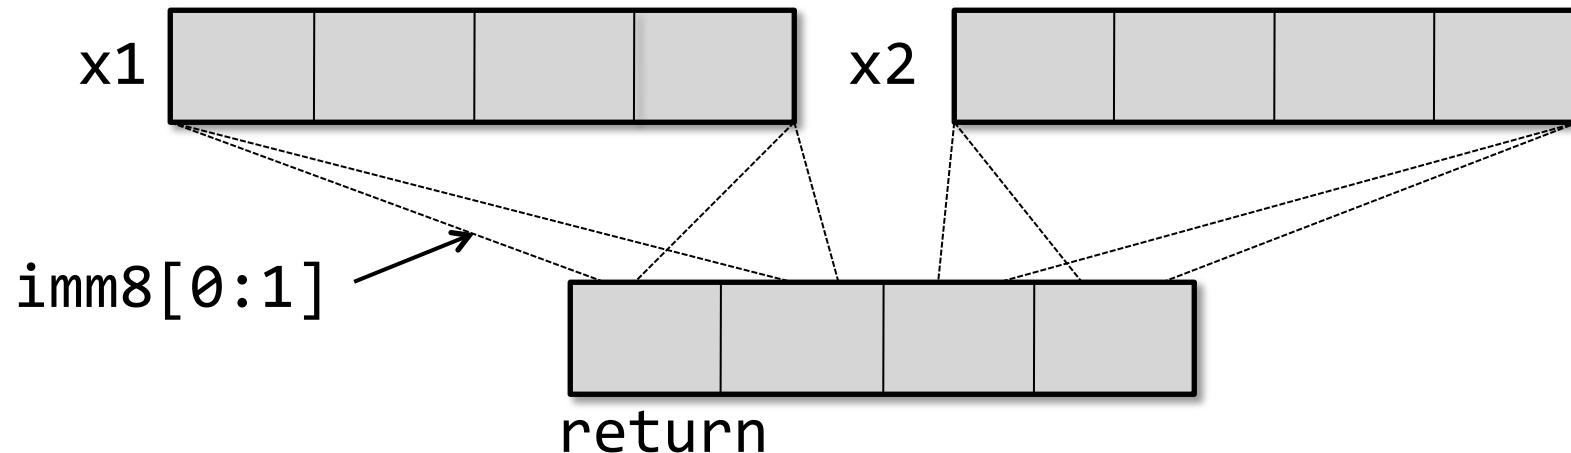
```
int[16] transpose(int[16] M) {  
    int[16] T = 0;  
    for (int i = 0; i < 4; i++)  
        for (int j = 0; j < 4; j++)  
            T[4 * i + j] = M[4 * j + i];  
    return T;  
}
```

This example comes from a Sketch grad-student contest

Implementation idea: parallelize with SIMD

Intel SHUFP (shuffle parallel scalars) SIMD instruction:

```
return = shufps(x1, x2, imm8 :: bitvector8)
```



High-level insight of the algorithm designer

Matrix M transposed in two shuffle phases

Phase 1: shuffle M into an intermediate matrix S with some number of shufps instructions

Phase 2: shuffle S into an result matrix T with some number of shufps instructions

Synthesis with partial programs helps one to complete their insight. Or prove it wrong.

The SIMD matrix transpose, sketched

```
int[16] trans_sse(int[16] M) implements trans {  
    int[16] S = 0, T = 0;
```

```
S[??::4] = shufps(M[??::4], M[??::4], ??);  
S[??::4] = shufps(M[??::4], M[??::4], ??);  
...  
S[??::4] = shufps(M[??::4], M[??::4], ??);
```



Phase 1

```
T[??::4] = shufps(S[??::4], S[??::4], ??);  
T[??::4] = shufps(S[??::4], S[??::4], ??);  
...  
T[??::4] = shufps(S[??::4], S[??::4], ??);
```



Phase 2

```
return T;  
}
```

The SIMD matrix transpose, sketched

```
int[16] trans_sse(int[16] M) implements trans {
    int[16] S = 0, T = 0;
    repeat (??) S[??::4] = shufps(M[??::4], M[??::4], ??);
    repeat (??) T[??::4] = shufps(S[??::4], S[??::4], ??);
    return T;
}

int[16] trans_sse(int[16] M) implements trans { // synthesized code
    S[4::4] = shufps(M[6::4], M[2::4], 11001000b);
    S[0::4] = shufps(M[11::4], M[6::4], 10010110b);
    S[12::4] = shufps(M[0::4], M[2::4], 10001101b);
    S[8::4] = shufps(M[8::4], M[12::4], 11010111b);
    T[4::4] = shufps(S[11::4], S[1::4], 10111100b);
    T[12::4] = shufps(S[3 From the contestant email:
    T[8::4] = shufps(S[4 Over the summer, I spent about 1/2
    T[0::4] = shufps(S[1 a day manually figuring it out.

}

```

Synthesis time: <5 minutes.

Demo: transpose on Sketch

Try Sketch online at <http://bit.ly/sketch-language>

The mechanics of program synthesis (constraint solving)

What to do with a program as a formula?

Assume a formula $S_P(x,y)$ which holds iff program $P(x)$ outputs value y

program: `f(x) { return x + x }`

formula: $Sf(x,y): y=x+x$

This formula is created as in program verification with concrete semantics [CMBC, Java Pathfinder, ...]

With program as a formula, solver is versatile

Solver as an **interpreter**: given x , evaluate $f(x)$

$$S(x,y) \wedge x=3 \quad \text{solve for } y \quad y \mapsto 6$$

Solver as a program **inverter**: given $f(x)$, find x

$$S(x,y) \wedge y=6 \quad \text{solve for } x \quad x \mapsto 3$$

This solver “bidirectionality” enables synthesis

Search of candidates as constraint solving

$SP(x,h,y)$ holds iff sketch $P[h](x)$ outputs y .

`spec(x) { return x + x }`

`sketch(x) { return x << ?? }` $S_{sketch}(x,y,h): y=x*2^h$

The solver computes h , thus synthesizing a program correct for the given x (here, $x=2$)

$S_{sketch}(x,y,h) \wedge x=2 \wedge y=4$ *solve for h* $\mathbf{h \mapsto 1}$

Sometimes h must be constrained on several inputs

$S(x_1,y_1,h) \wedge x_1=0 \wedge y_1=0 \wedge$

$S(x_2,y_2,h) \wedge x_2=3 \wedge y_2=6$ *solve for h* $\mathbf{h \mapsto 1}$

Inductive synthesis

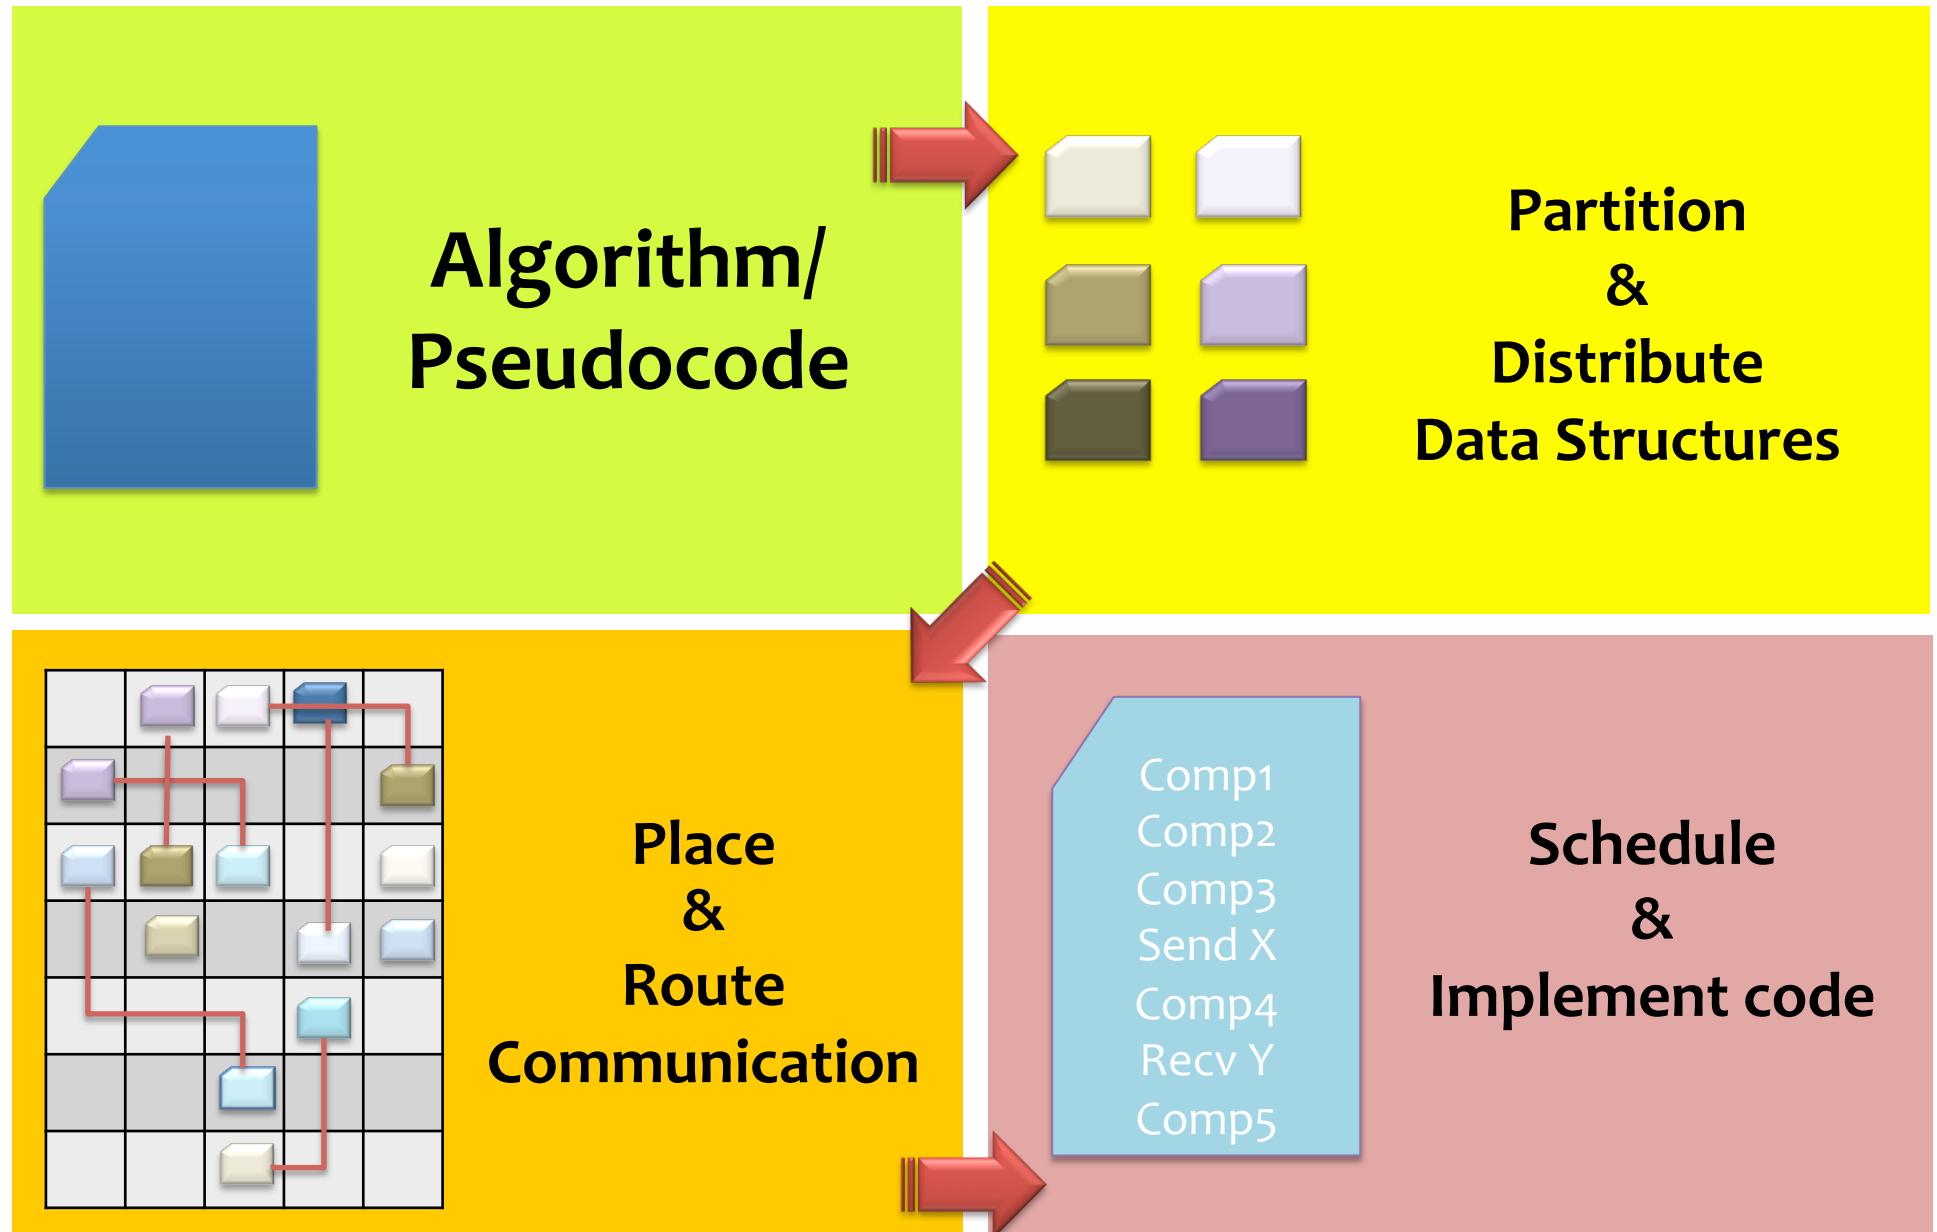
Our constraints encode **inductive synthesis**:

We ask for a program P correct on a few inputs.

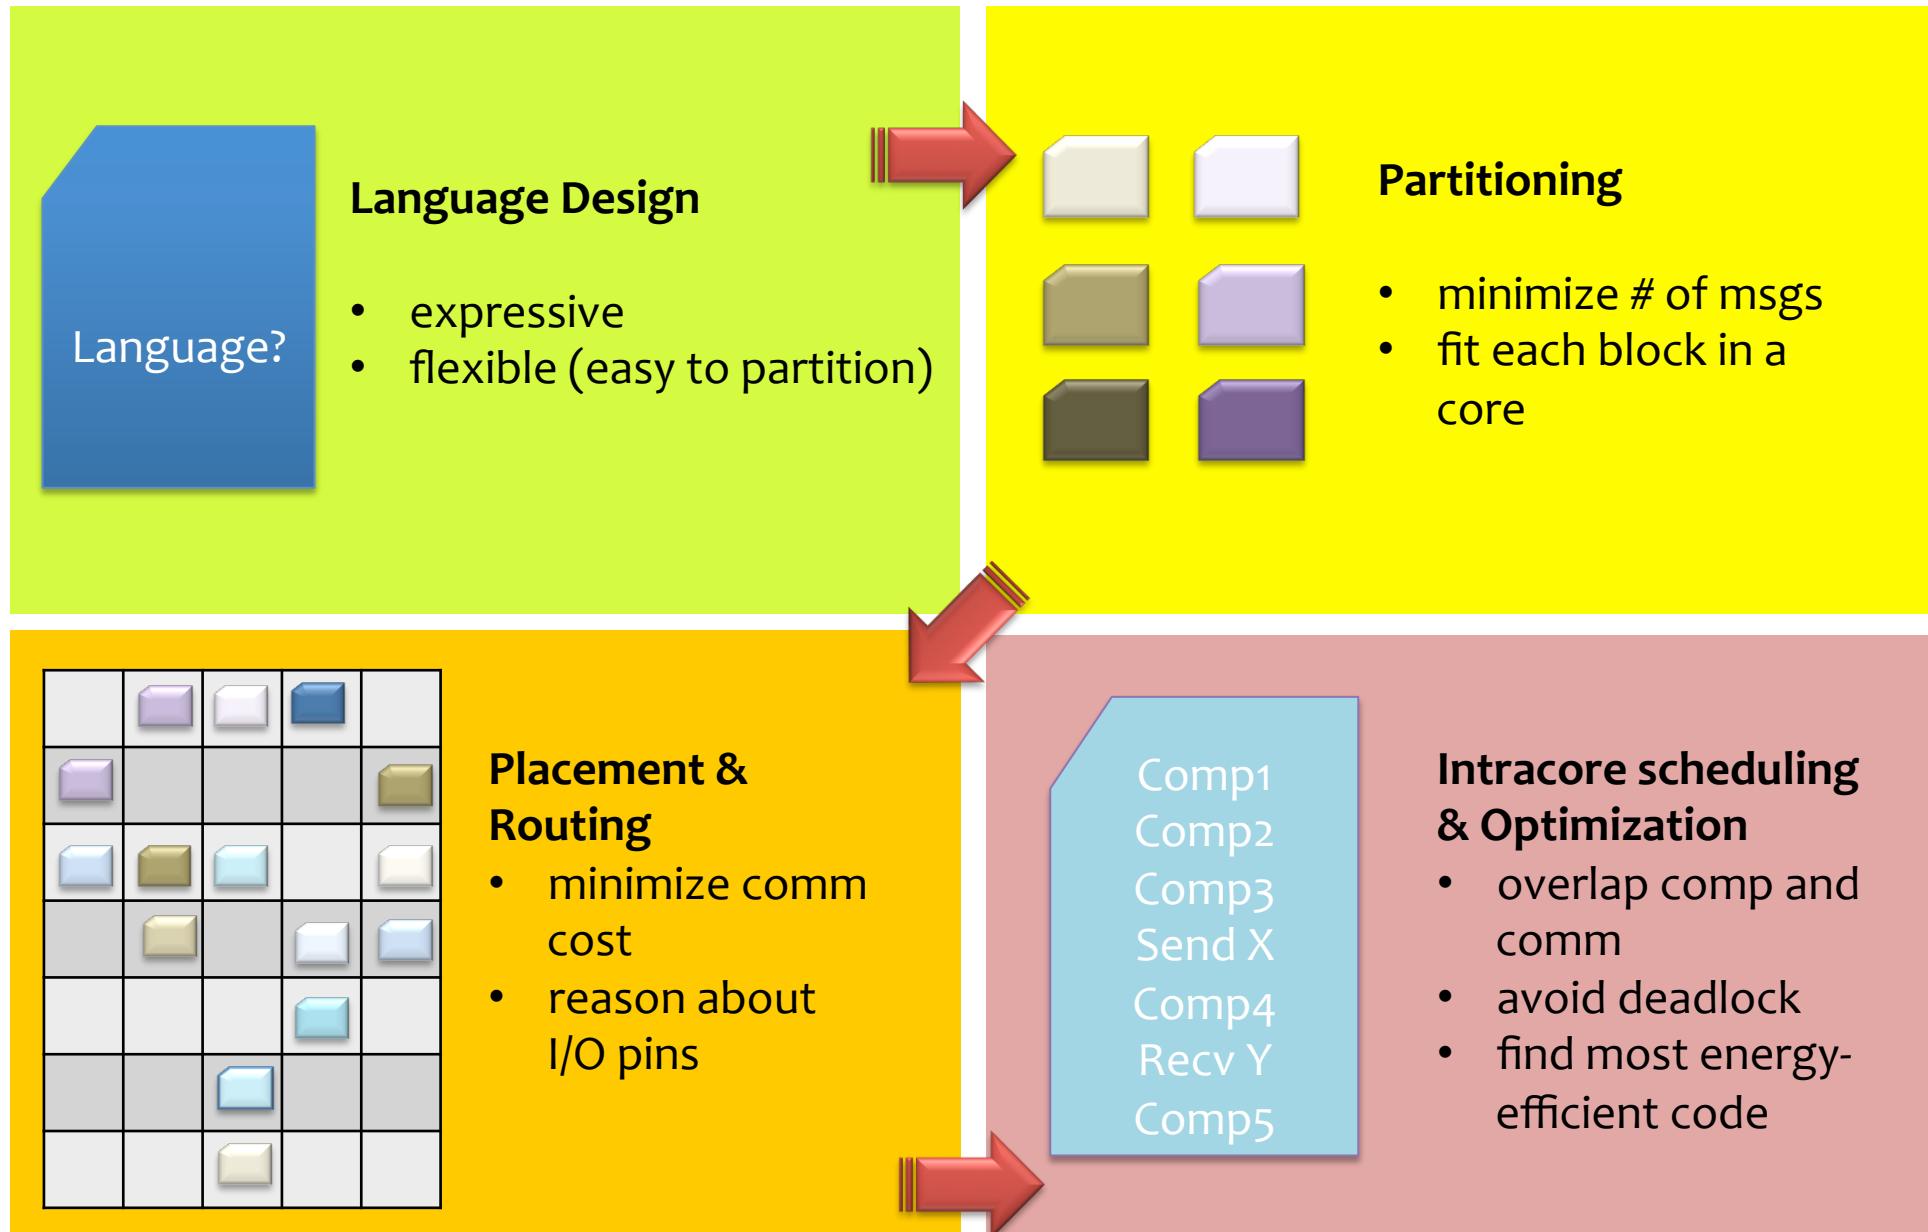
We hope (or test, verify) that P is correct on rest of inputs.

We propose a programming model
for low-power devices by
exploiting program synthesis.

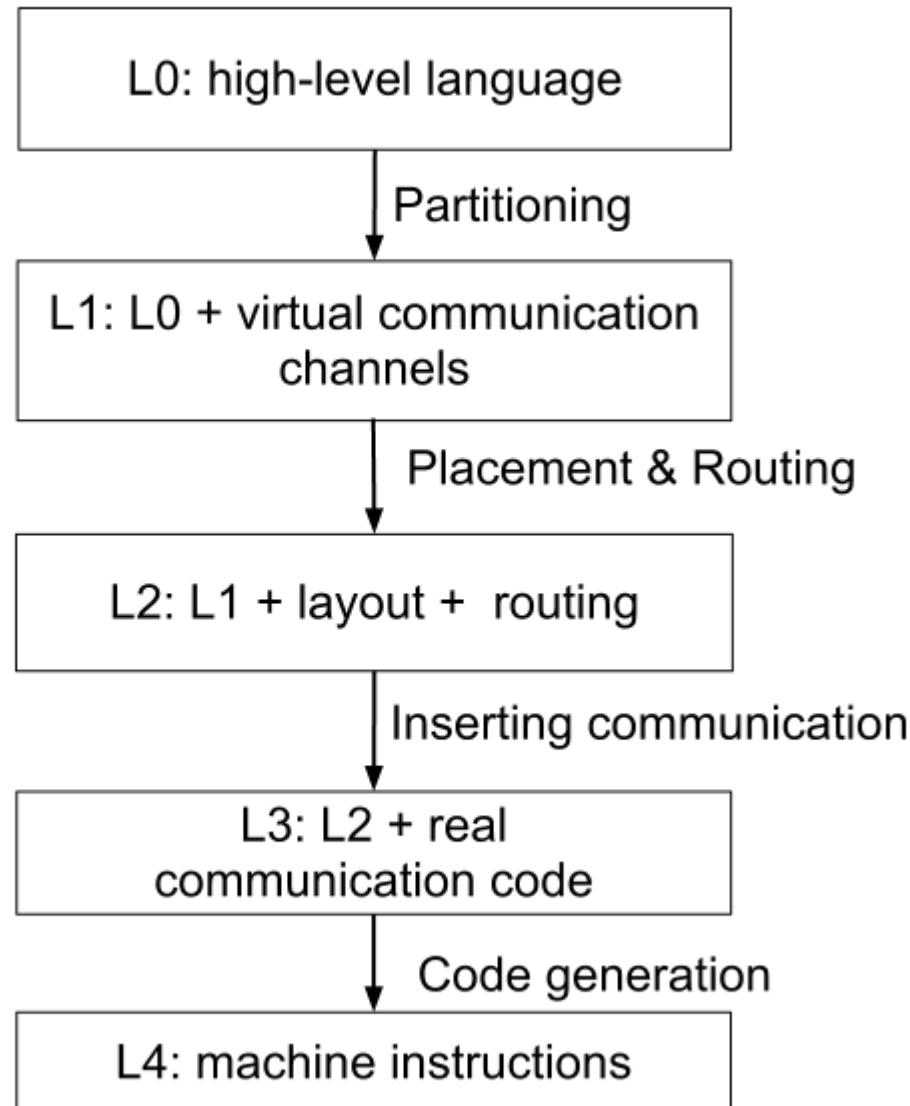
Programming/Compiling Challenges



Project Pipeline

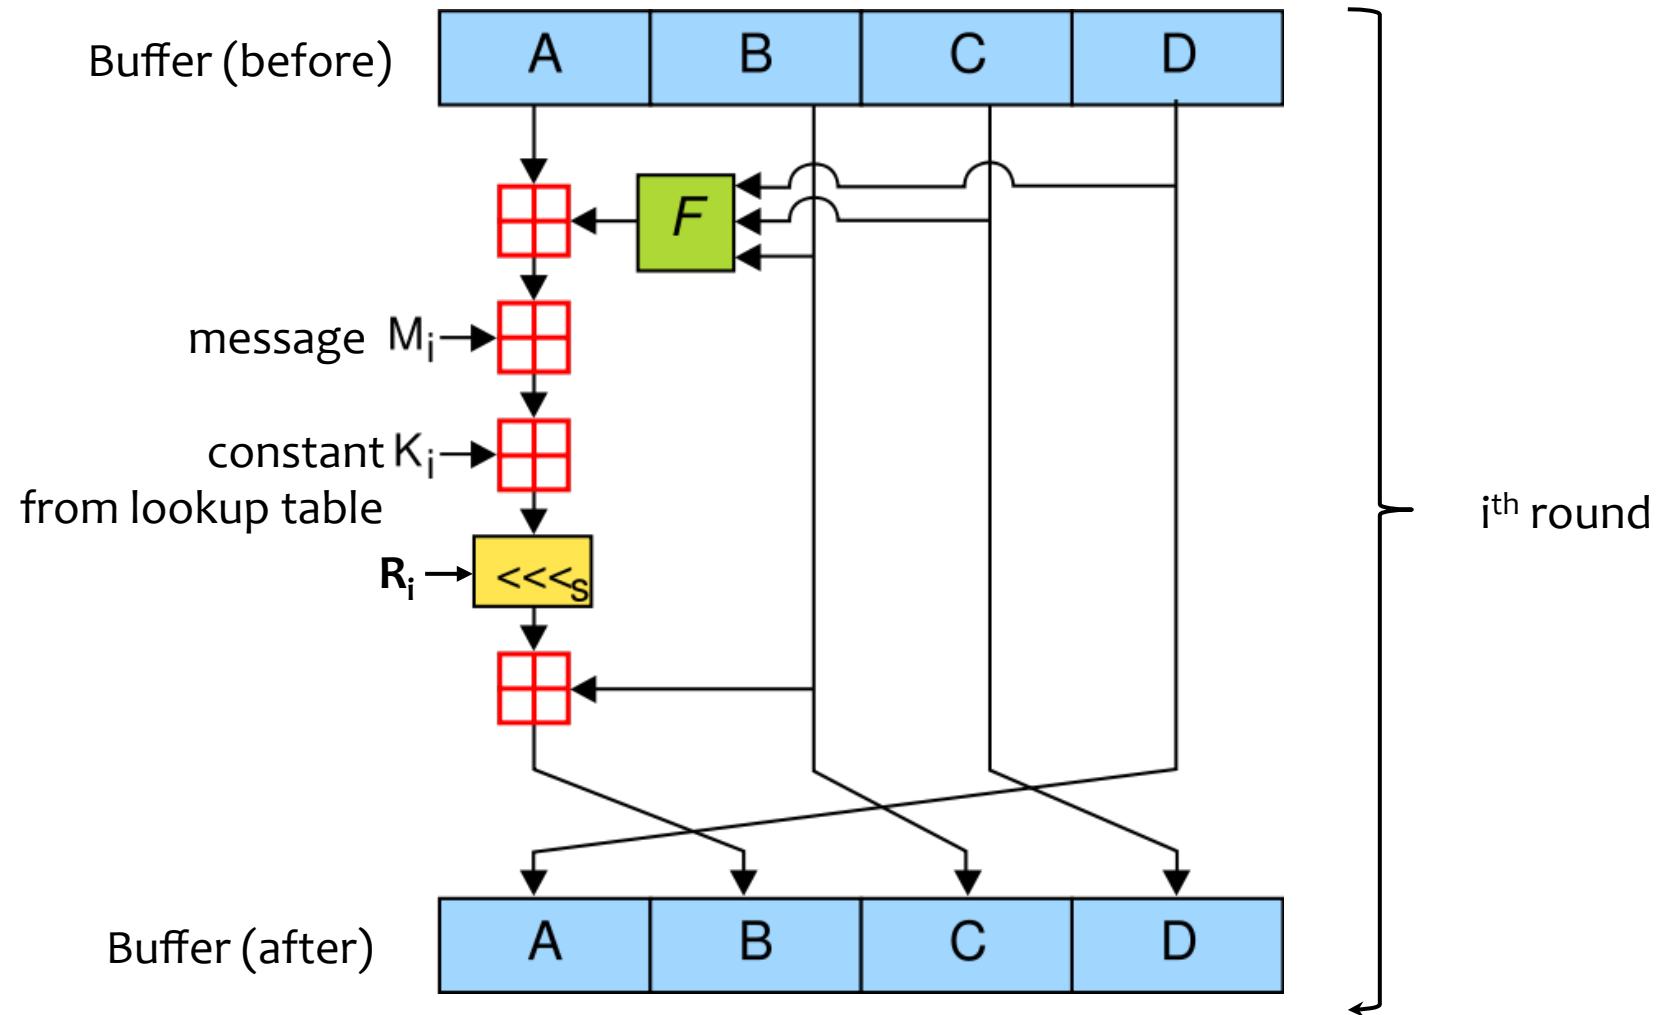


Programming model abstractions



MD5 case study

MD5 Hash



MD5 from Wikipedia

```
//Process the message in successive 512-bit chunks:  
for each 512-bit chunk of message  
    break chunk into sixteen 32-bit words w[j], 0 ≤ j ≤ 15  
//Initialize hash value for this chunk:  
var int a := h0  
var int b := h1  
var int c := h2  
var int d := h3  
//Main loop:  
    for i from 0 to 63  
        if 0 ≤ i ≤ 15 then  
            f := (b and c) or ((not b) and d)  
            g := i  
        else if 16 ≤ i ≤ 31  
            f := (d and b) or ((not d) and c)  
            g := (5×i + 1) mod 16  
        else if 32 ≤ i ≤ 47  
            f := b xor c xor d  
            g := (3×i + 5) mod 16  
        else if 48 ≤ i ≤ 63  
            f := c xor (b or (not d))  
            g := (7×i) mod 16  
        temp := d  
        d := c  
        c := b  
        b := b + leftrotate((a + f + k[i] + w[g]) , r[i])  
        a := temp  
    end for  
//Add this chunk's hash to result so far:  
h0 := h0 + a  
h1 := h1 + b  
h2 := h2 + c  
h3 := h3 + d  
end for  
  
var char digest[16] := h0 append h1 append h2 append h3 // (Output is in little-endian)  
  
//leftrotate function definition  
leftrotate (x, c)  
    return (x << c) binary or (x >> (32-c));
```

Actual MD5 Implementation on GA144

```
900 List
md5 host offset 508 0 ;
relocation helpers 32 fh load target
data
 106 +n high 4 fh load 40
  6 +n low  6 fh load 40
const and adder and rotator
 105 +n high 8 fh load 12 fh load 27
  5 +n low 10 fh load 12 fh load 36
md5
 104 +n high 14 fh load 40
  4 +n low 14 fh load 40
msg
 103 +n high 0 org 16 fh load 39
  3 +n low 0 org 16 fh load 39
rots
 102 +n high 18 fh load 3B
  2 +n low 20 fh load 3B
entry point
 101 +n high 22 fh load 1A
  1 +n low 24 fh load 40
octet feeder
  0 +n 26 fh load 40
 100 +n 28 fh load 35
 200 +n 30 fh load 3F
```

```
902 List
- load descriptor offset 508 0 ;
inward offset 100 / 1 and drop
  if up ; then down ;
outward offset 100 / 1 and drop
  if down ; then up ;
west com @ 1 and drop if right ; then left ;
/node n offset + dup +node /ram ;

106 /node west /p 6 /node west /p
105 /node west /p 5 /node west /p
104 /node west /p 4 /node west /p
103 /node west /p 3 /node west /p
102 /node west /p 2 /node west /p
101 /node inward /p 1 /node west /p
  0 /node inward /p 100 /node outward /p
200 /node 2A /p

+break slot addr node offset + break ;
0 34 100 +break 3 3F 0 +break
```

```
932 List
- relocating code
+n offset + node ;
inward offset 100 / 1 and drop
  if ' ---u lit ; then ' -d-- lit ;
outward offset 100 / 1 and drop
  if ' -d-- lit ; then ' ---u lit ;
col com @ 1 and drop ;
east col if ' --l- lit ; then ' r--- lit ;
west col if ' r--- lit ; then ' --l- lit ;
greatest col if r--- ; then --l- ;
greatest col if --l- ; then r--- ;
```

```
908 List
- 105 high constant generator and adder
14 org
rotate 14 ni-n inward b! dup dup or
a! push push a pop pop
-if - push push !b pop !b pop then
for +* unext !b drop !b
a 2/ 2/ FFFF and or ;
+tc 22 nn-n + inward b! !b . + FFFF and ;
```

```
910 List
- 5 low constant generator rotator and adder
4 org
rotate 14 ni-n inward b! dup dup or
a! push push a pop pop
-if - push !b !b pop then
for +* unext !b push !b drop pop
a 2/ 2/ FFFF and or ;
+tc 22 nn-n + inward b! dup -if 2* -if
  3 !b drop FFFF and ; then
  2 !b drop FFFF and ; then
2* -if 1 !b drop FFFF and ; then
dup or !b FFFF and ;
```

```
912 List
- 105 5 constant generator rotator adder
bare * 0 org
st 00 -n east b! !p .. !b !b ;
> 04 west a! east b!
> !b .. dup or a! ..
3 for b @ f @ m @ . + t get +c
push s @ rotate +c pop a! a! next ;
/* org
```

```
4 List
104 4 md5 buffer code 0 org
  , 0 , 0 ,
  ep 04 -3 west b! 3 dup dup or a! ;
iss 08 prep for +!b unext ;
ab 0B prep for !b +!unext ;
ip 0B a 3 and a! ;
  11 xyz-n push over - push and pop pop and
or 14 nn-n over FFFF or and or ;
  17 xyz-n a! push a and pop a - and +or ;
  1B xyz-n or or ;
  1C xyz-n a! push a - +or pop or ;
  ss 1F west b! !b send east b! !b ;
round 23 2E a! ! dup dup or a!
15 for a + clip b + clip
dup b send c + clip d + clip
a push .. 2E f' . + pop a!
f send m pass s pass a @b 34/1 !+
@+ drop +! drop clip next ;
md5 3B
@p round f' @p round g'
@p round h' @p round ; i'
```

```
916 List
- 103 3 message buffer code 0 org
0 , 0 , 0 , 0 , 0 , 0 , 0 , 0 ,
0 , 0 , 0 , 0 , 0 , 0 , 0 , 0 ,
msgs 10 ni pop a! +!+ a push
15 for dup a! @ !b over . +
F and west a! @ !b next drop drop 1C/2 ;
msg 1C east b! !p .. !p b! !p ..
!b west !b !p .. go ..
!b !p !b !p .. !b grab .. md5 ..
west a! 3 dup push
```

```
for @ !b unext !b
msgs 1 , 0 , f's
msgs 5 , 1 , g's
msgs 3 , 5 , h's
msgs 7 , 0 , i's
@p !b .. toss ..
begin !b ! unext ;
```

```
918 List
- 102 rots generator 0 org
29E38 , B16A , 28F4B , A059 ,
send n dup east a! ! inward a! ! ;
keep 09 inward b! east a! 3 for @ !b unext ;
put OF inward b! east a! 3 for !b ! unext ;
ncol 15 for 2/ unext
0col 17 F and ; lcol 19 3 ncol ;
2col 1B dup push 7 ncol pop
-if drop - dup then drop ;
3col 21 11 ncol - ;
jump 24 i pop + push ;
rots 25 i dup 2/ 2/ 2/ 2/ b! !b over
3 and jump 0col ; lcol ; 2col ; 3col ;
rotgen 2E east a! !p .. msg .. ! put
0 63 for dup rots send 1 . + next
drop keep ;
```

```
920 List
- 2 save and add abcd back 0 org
0 , 0 , 0 , 0 ,
0 , 0 , 0 , 0 ,
put 08 dup dup or a! 3 for
east b! +! b inward b! +! b next ;
sum 12 dup dup or a! 3 for
east b! !b @ . + dup FFFF and !+
inward b! !b @ . + over 2*
-if drop 1 . + dup then drop
FFFF and !+ next ;
sums 25 east a! !p .. msg .. 28 ! put
east a! 63 for !b ! unext sum ;
send 2B west b! dup dup or a!
7 for @+ !b unext ;
/buf 35 dup dup or a! west b!
7 for @b !+ unext 3B/3 ; 3B
```

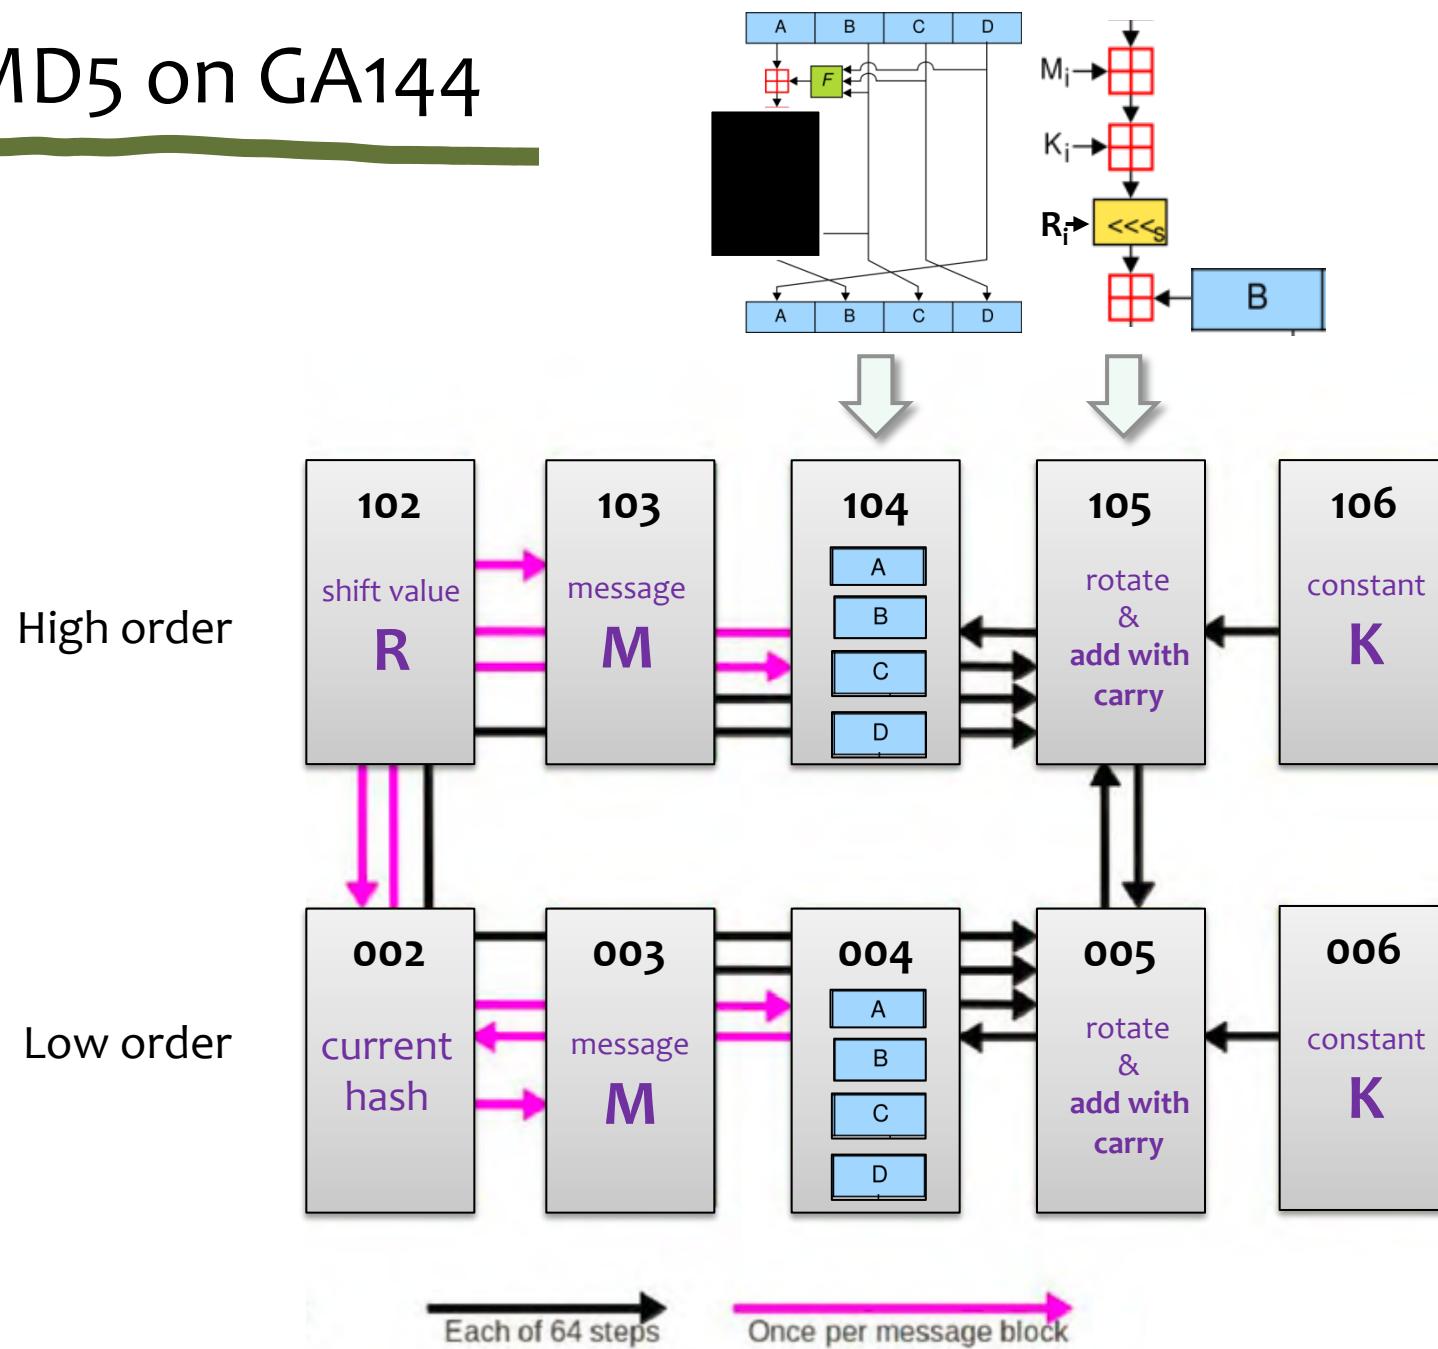
```
922 List
- 101 main control 0 org
code 102 west b! 20 for @p !b unext ..
103 dup dup or a! 15 for
@p !+ unext .. 09 gowest ; goeast ; OB
highs 0B east b! dup dup or a!
8 for @+ !b unext a push inward a!
15 for @ !b unext pop a!
15 for @ !b unext pop a!
+!b !b .. !b .. !p !b ; .. rotgen ..
```

```
924 List
- 1 main control 0 org
code 2 west b! 20 for @p !b unext ..
3 dup dup or a! 15 for
@p !+ unext .. 09 gowest ; goeast ; OB
half 0B p push west a! @ @
7 for 2* unext or pop a! !
lows 12 inward b! !p .. highs .. !b
east b! dup dup or a!
8 for @+ !b unext a push
15 for east half inward half next
pop a! +! !b +! !b .. !p !b ; .. sums ..
pass east b! west .. a! !p !b .. send ..
7 for @b ! unext ;
prime 2F east b! 38 a! !p .. /buf ..
!b 7 for @+ !b unext 37/3 gowest ..
2301 , 6745 , AB89 , EFCD ,
DCFE , 98BA , 5476 , 1032 , 40
```

```
926 List
- 0 octet feeder 0 org 0 , 0 , 0 ,
add 03 n-n @ . + dup FFFF and !+
count 06 1 a! 8 add 2* -if
drop 1 add then drop ;
get OF @ -if dup or a! 80 ; then
a! if count then a! ;
octets 16 n for get !b next ;
ablk 1A n 55 octets
a if drop 7 octets ; then drop
@+ drop a for
@+ dup FF and !b 7 for 2/ unext
FF and !b next
0 dup dup or a! 3 for @ !b unext warm ;
start 30 east b! inward a!
begin @p !b .. lows .. ablk end
digest east b! @p .. pass .. !b inward a!
7 for @b ! unext 40/3 ;
```

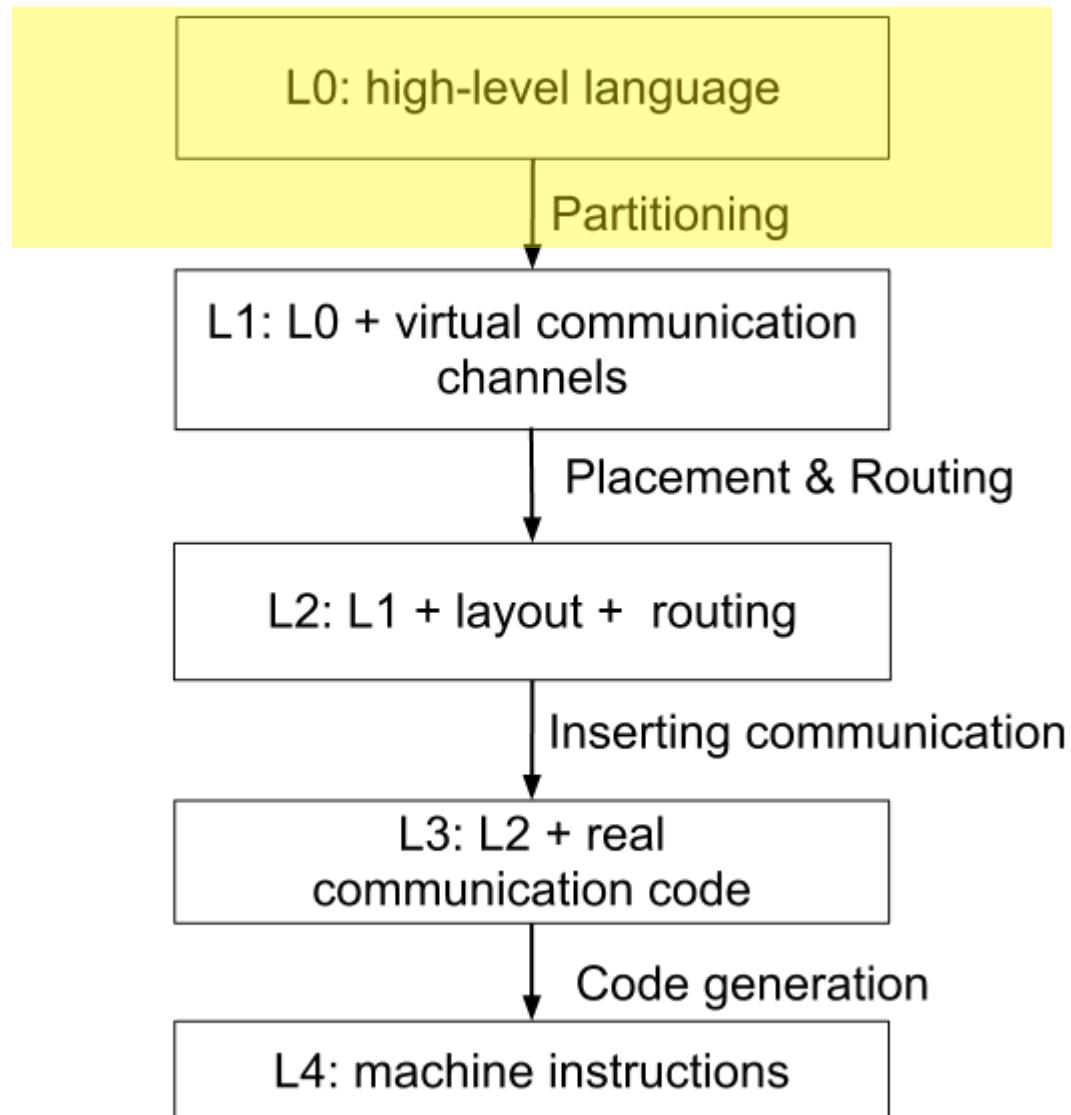
```
928 List
- 100 application interface
addresses 0-7 reserved for result 8 org
/count 08 inward b! 0 1 ..
@p !b !b dup .. @p a! @p dup ..
!b @p .. !+ !+ ..
/digest 10 8 a! 3 for @p !b unext ..
  east b! .. @p !b .. prime ..
/msg @p !b .. start ..
up a! @ begin !b @ -until
!b 0 a! @p .. digest .. !b
7 for @b !+ unext 24/3 warm ; .. /count ; 25
```

MD5 on GA144



This is how we express MD5

Project Pipeline

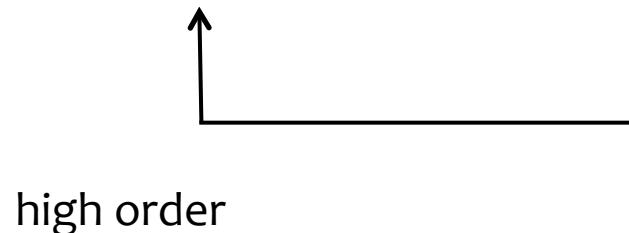
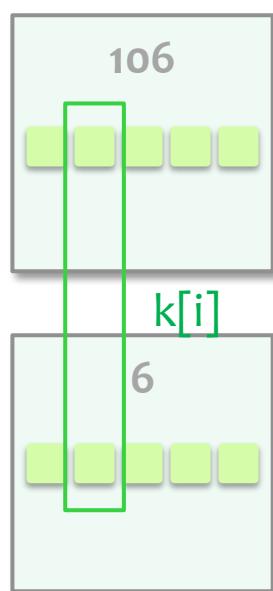


Annotation at Variable Declaration

```
typedef pair<int,int> myInt;
```

```
vector<myInt>
vector<myInt>
vector<myInt>
vector<myInt>
vector<int>
```

```
message[16];
k[64];
output[4];
hash[4];
r[64];
```



@core indicates where data lives.

Annotation in Program

(104,4) is home
for md5() function

```
void@(104,4) md5() {  
    for(myInt@any t = 0; t < 16; t++) {  
        myInt@here a = hash[0], b = hash[1], c = hash[2], d = hash[3];  
        for(myInt@any i = 0; i < 64; i++) {  
            myInt@here temp = d;  
            d = c;  
            c = b;  
            b = round(a, b, c, d, i);  
            a = temp;  
        }  
        hash[0] += a;  
        hash[1] += b;  
        hash[2] += c;  
        hash[3] += d;  
    }  
    output[0] = hash[0];  
    output[1] = hash[1];  
    output[2] = hash[2];  
    output[3] = hash[3];  
}
```

@any suggests that any core can have this variable.

@here refers to the function's home @ (104,4)

MD5 in CPart

```
typedef pair<int,int> myInt;

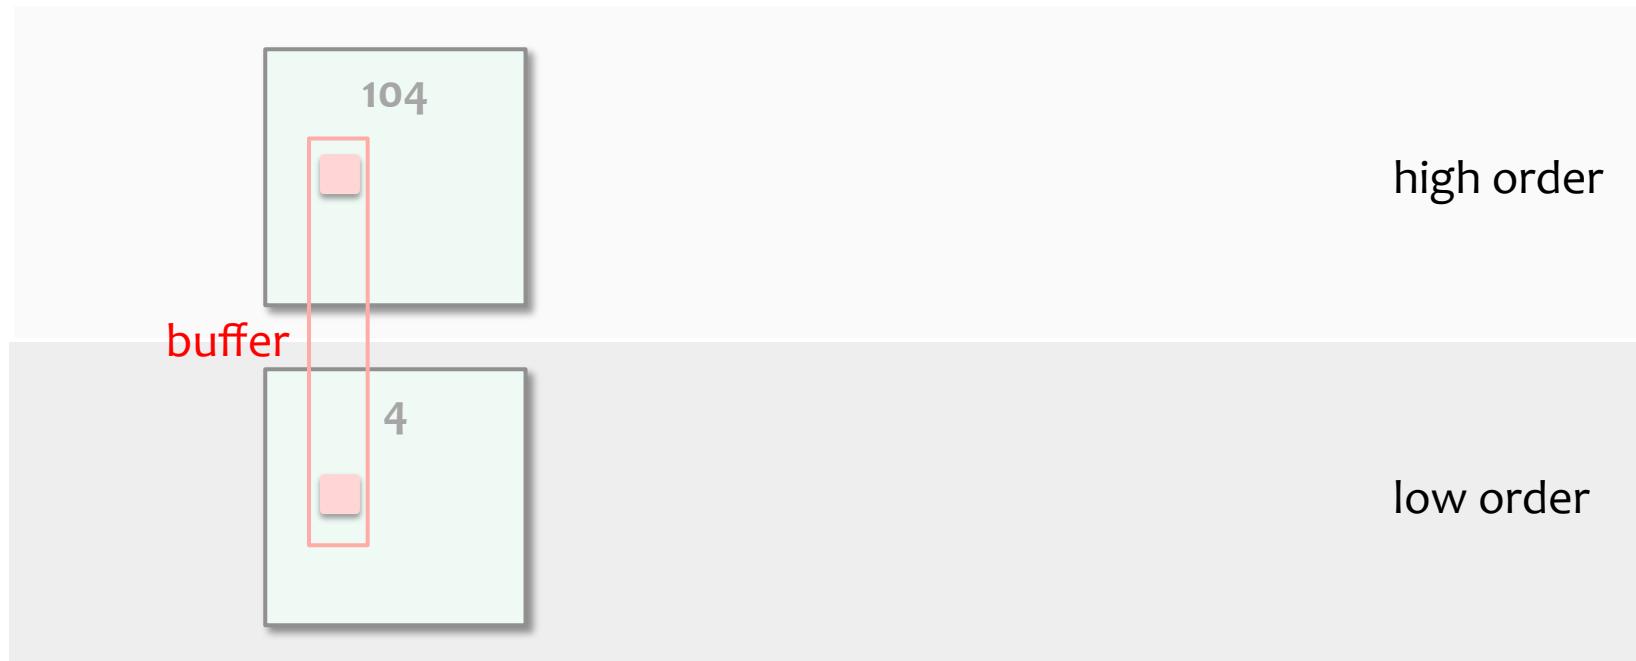
vector<myInt>@[ 0:64]=(106,6) k[64];

myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {
    myInt@h sum = buffer +@here k[i] + message[g];
    ...
}
```

MD5 in CPart

```
typedef pair<int,int> myInt;  
  
vector<myInt>@[ 0:64]=(106,6) } k[64];  
  
myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {  
    myInt@h sum = buffer +@here k[i] + message[g];  
    ...  
}
```

buffer is at (104,4)



MD5 in CPart

```
typedef pair<int,int> myInt;  
  
vector<myInt>@[0:64]=(106,6) k[64];
```

```
myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {  
    myInt@h sum = buffer +@here k[i] + message[g];  
    ...  
}
```

k[i] is at (106, 6)



MD5 in CPart

```
typedef pair<int,int> myInt;  
vector<myInt>@{[0:64]=(106,6)} k[64];
```

```
myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...){  
    myInt@h sum = buffer +@here k[i] + message[g];  
    ...  
}
```

+ is at (105,5)

105

+

high order

5

+

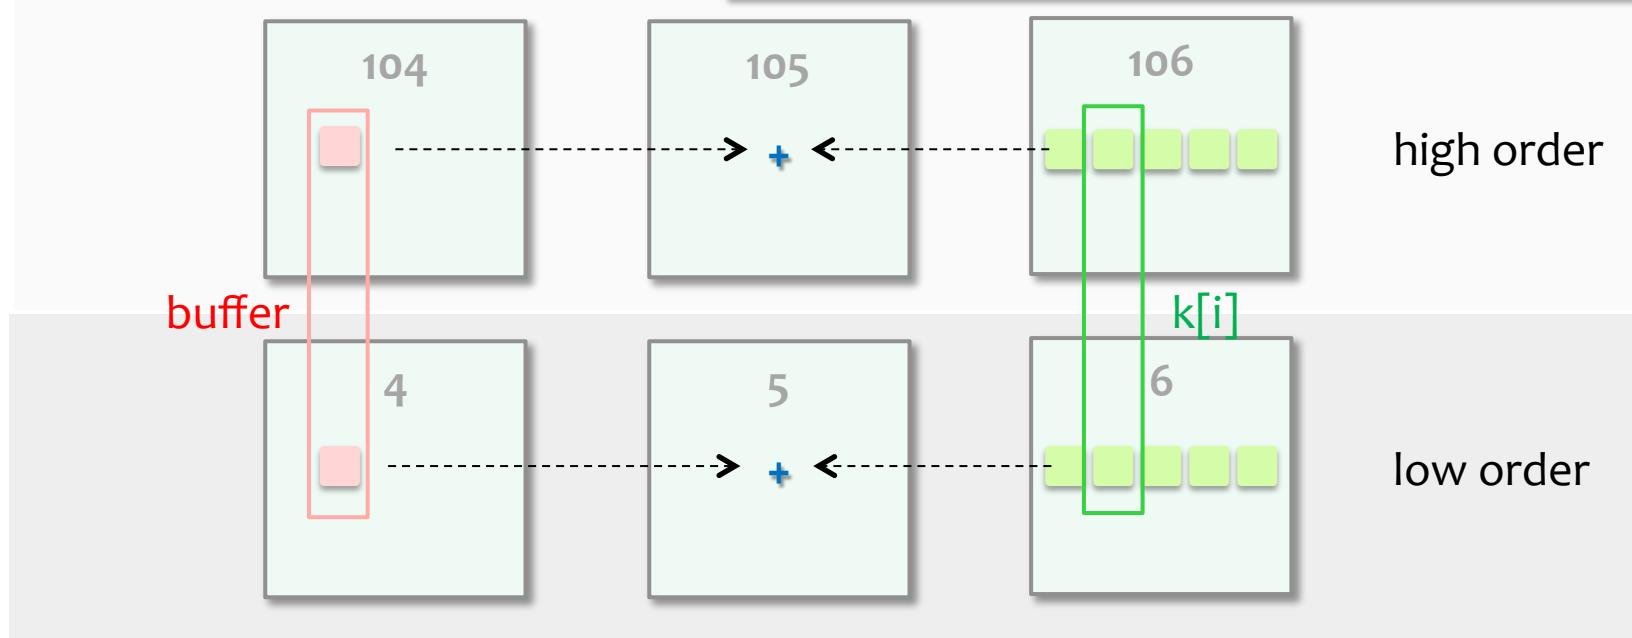
low order

MD5 in CPart

```
typedef pair<int,int> myInt;  
  
vector<myInt>@[0:64]=(106,6) k[64];  
  
myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {  
    myInt@h sum = buffer +@here k[i] + message[q];  
    ...  
}
```

buffer is at (104,4)
+ is at (105,5)
k[i] is at (106,6)

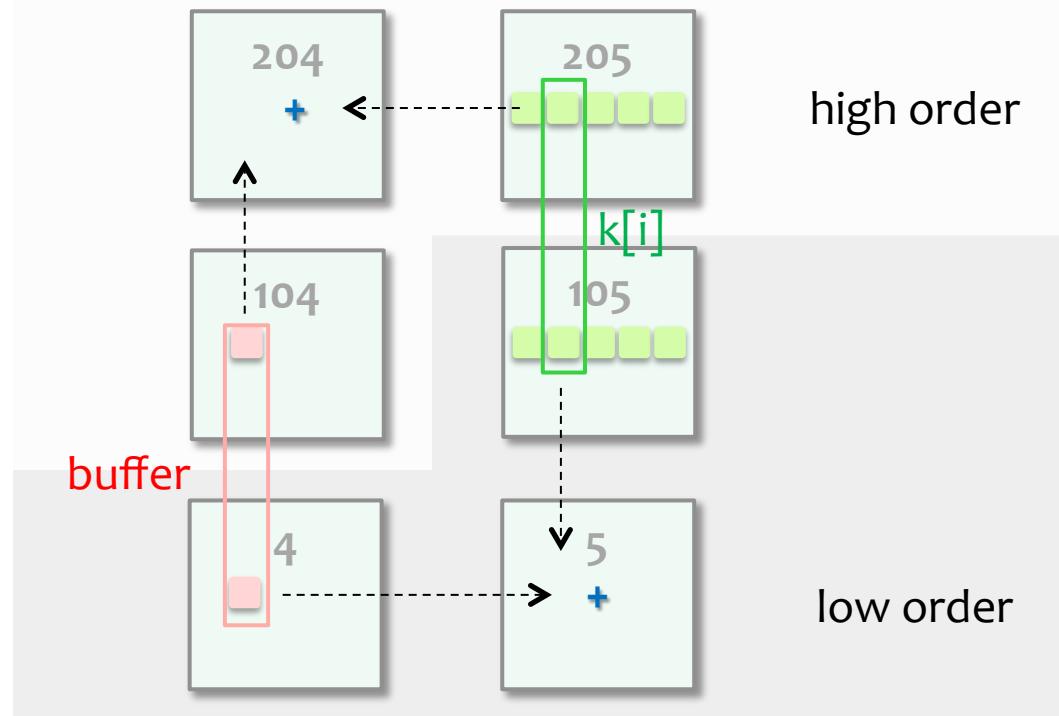
Implicit communication in source program.
Communication inserted by synthesizer.



MD5 in CPart

```
typedef pair<int,int> myInt;  
  
vector<myInt>@[0:64]=(205,105)} k[64];  
  
myInt@(204,5) sumrotate(myInt@(104,4) buffer, ...){  
    myInt@h sum = buffer +@here k[i] + message[g];  
    ...  
}
```

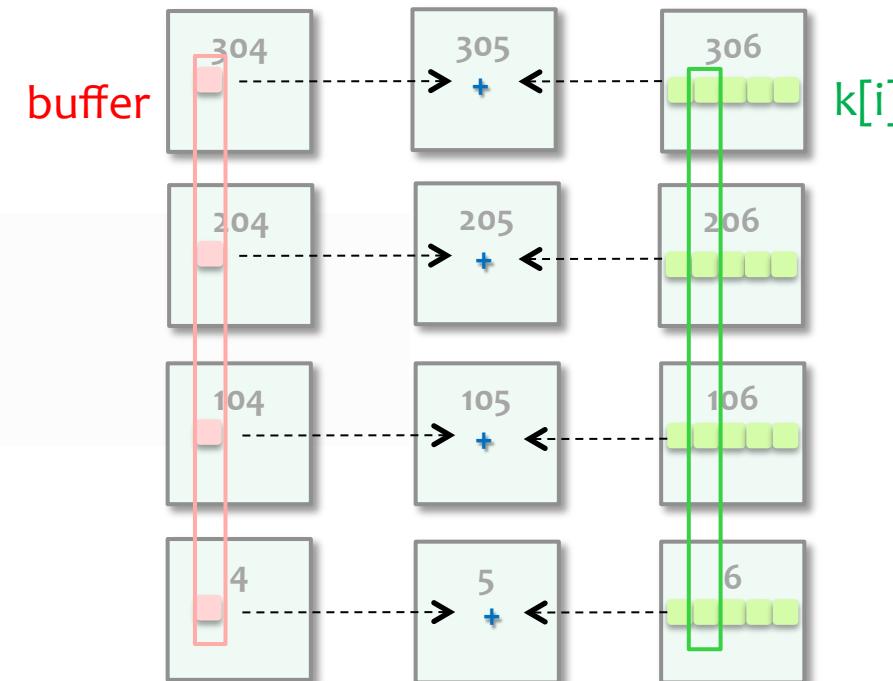
buffer is at (104,4)
+ is at (204,5)
k[i] is at (205,105)



MD5 in CPart

```
typedef vector<int> myInt;  
  
vector<myInt>@[0:64]={306,206,106,6} k[64];  
  
myInt@{305,205,105,5} sumrotate(myInt@{304,204,104,4} buffer, ...){  
    myInt@h sum = buffer +@here k[i] + message[g];  
    ...  
}
```

buffer is at {304,204,104,4}
+ is at {305,205,105,5}
k[i] is at {306,206,106,6}



MD5 in CPart

```
typedef pair<int,int> myInt;

vector<myInt>@[0:64]=(106,6) k[64];

myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {
    myInt@h sum = buffer +@here k[i] + message[g];
    ...
}
```

MD5 in CPart

```
typedef pair<int,int> myInt;  
  
vector<myInt>@[0:64]=(106,6) k[64];  
  
myInt@(105,5) sumrotate(myInt@(104,4) buffer, ...) {  
    myInt@h sum = buffer +@here k[i] +@?? message[g];  
    ...  
}
```



+ happens at **here**
which is **(105, 5)**

+ happens at where the
synthesizer decides

Summary: Language & Compiler

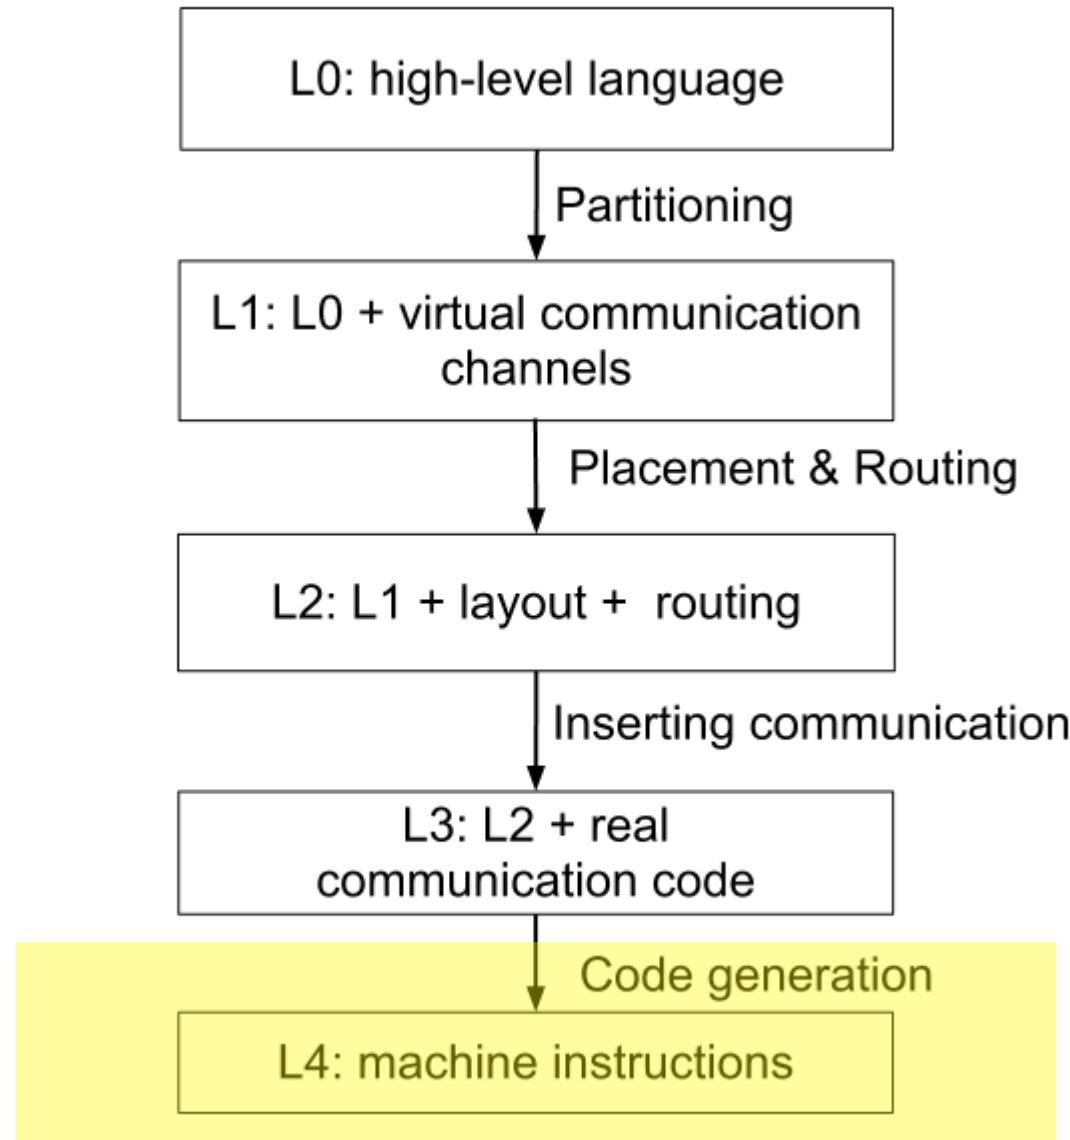
Language features:

- Specify code and data placement by annotation
- No explicit communication required
- Option to not specifying place

Synthesizer fills in holes such that:

- number of messages is minimized
- code can fit in each core

Project Pipeline



Program Synthesis for Code Generation

Synthesizing optimal code

Input: unoptimized code (the spec)

Search space of all programs

Synthesizing optimal library code

Input: sketch + spec

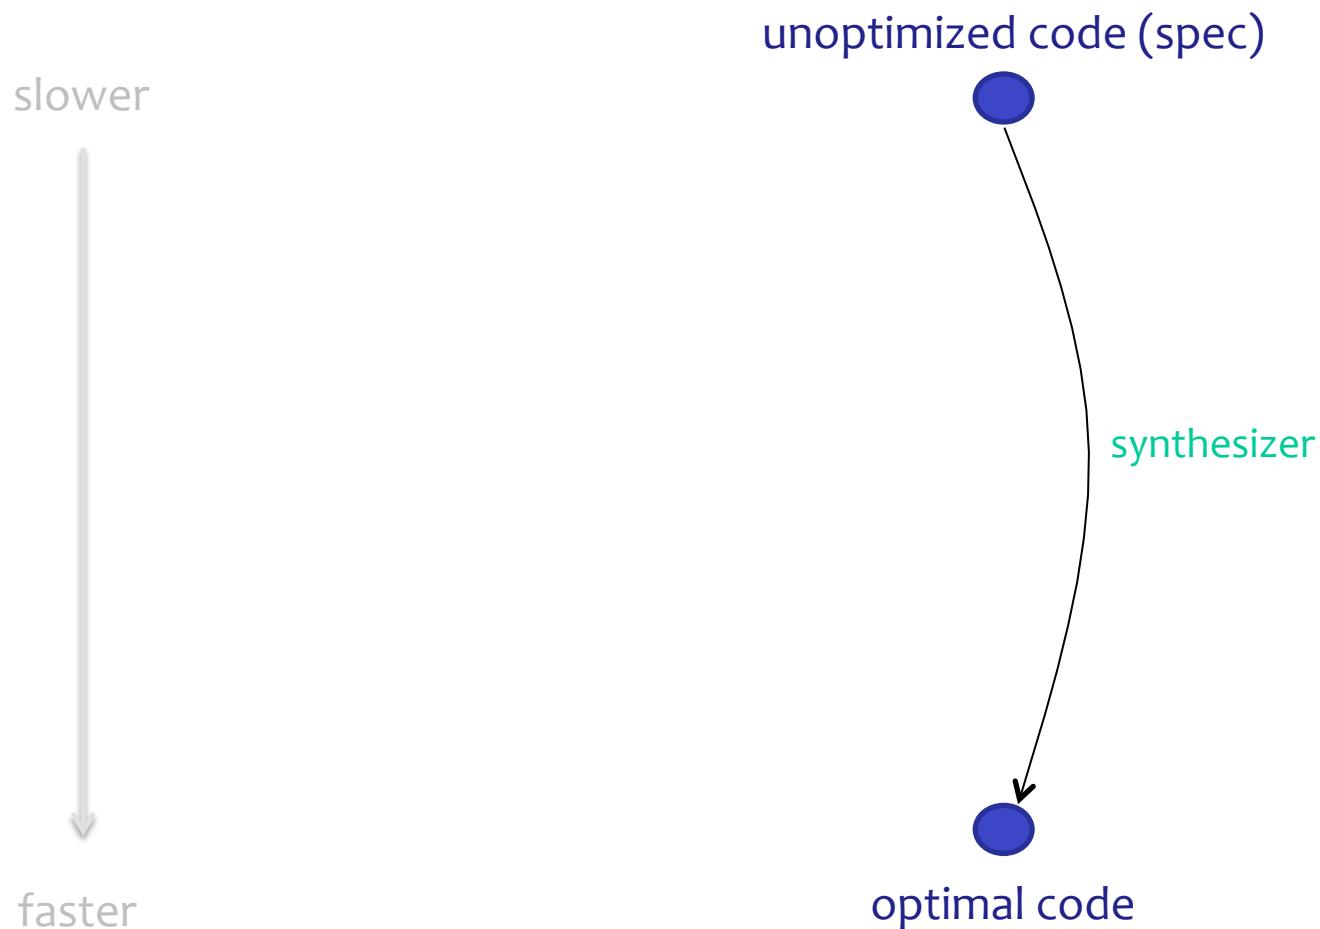
Search completions of the sketch

Synthesizing communication code

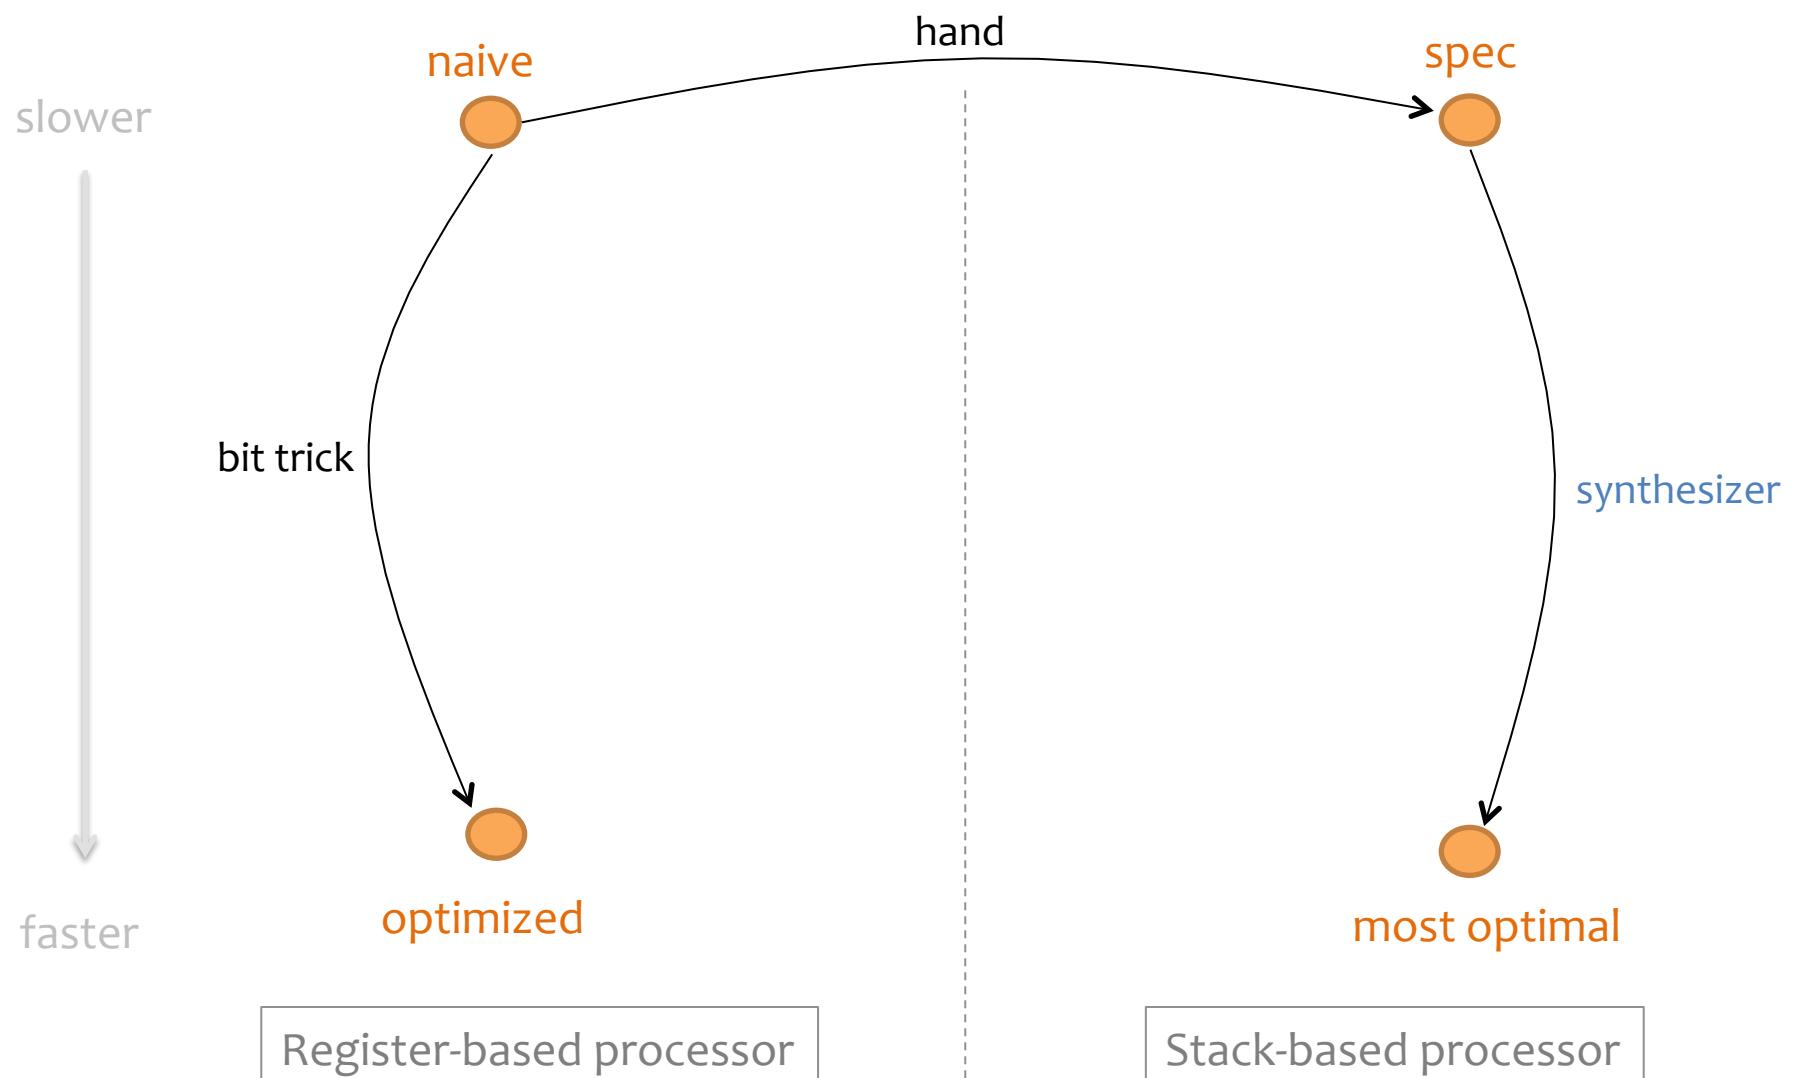
Input: program with virtual channels

Insert actual communication code

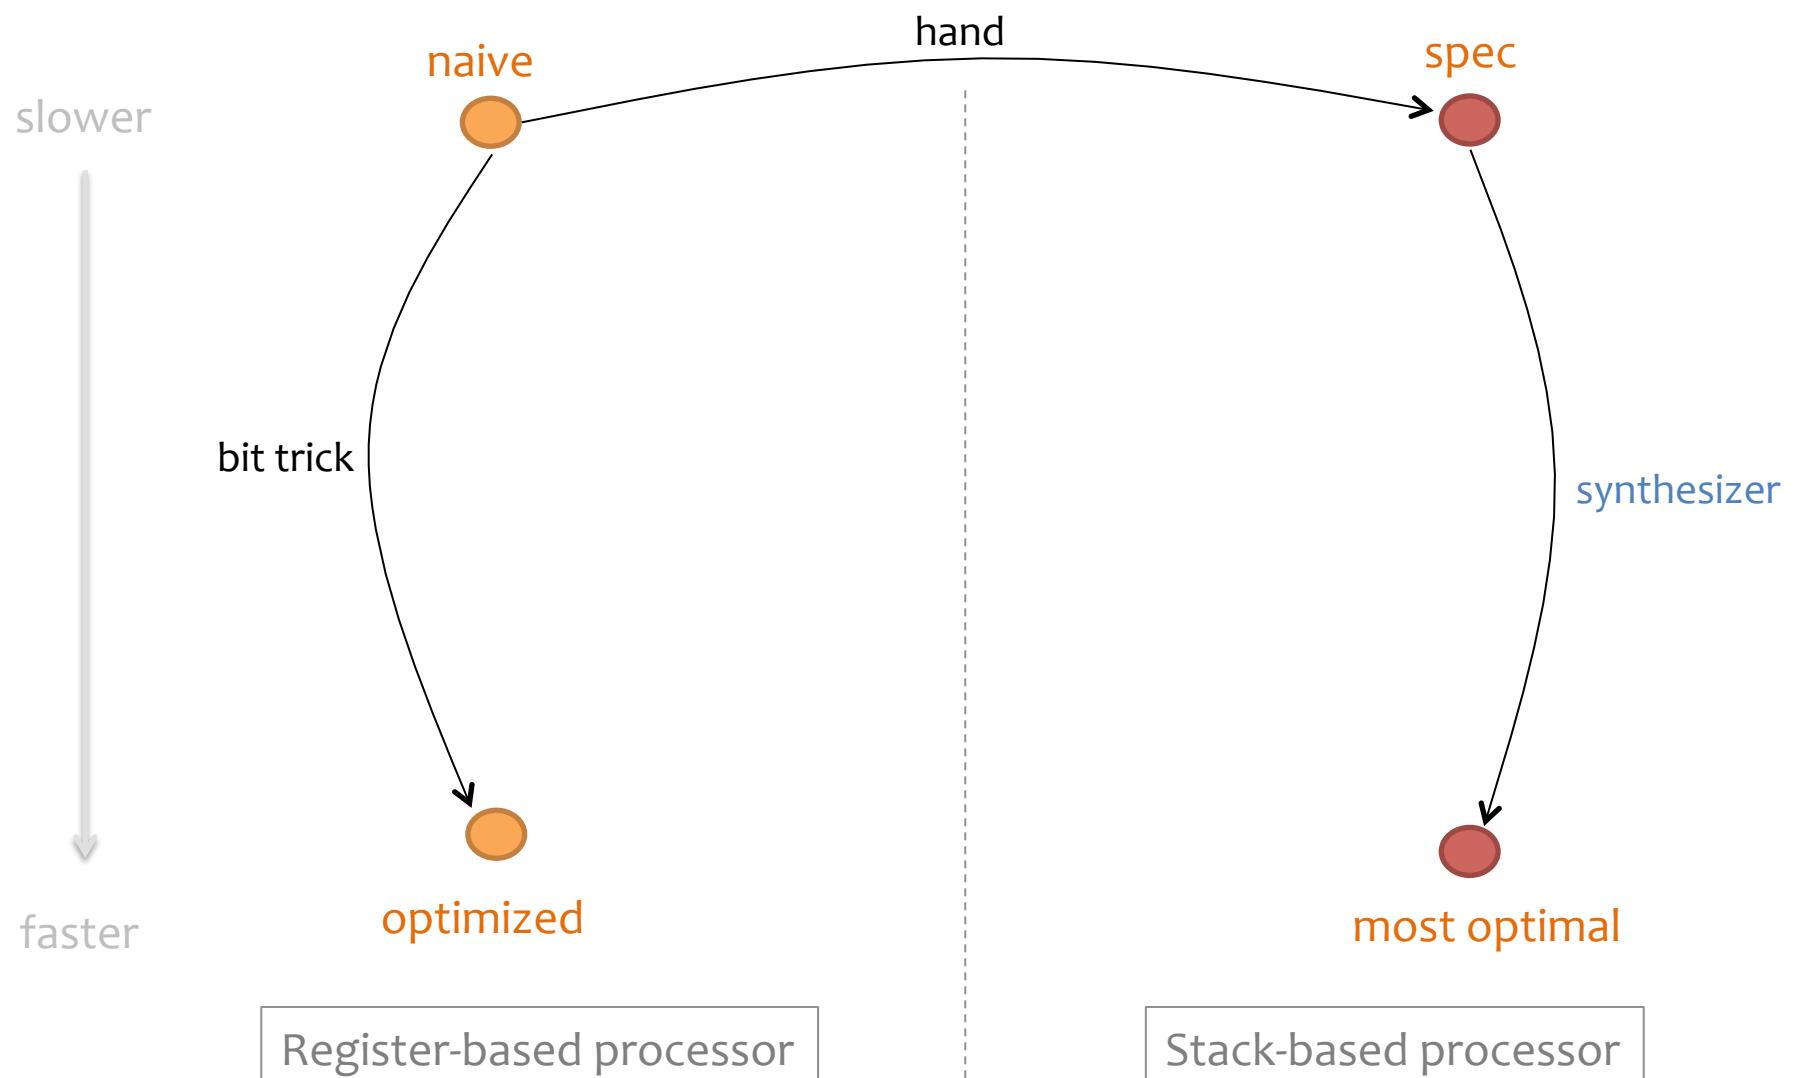
1) Synthesizing optimal code



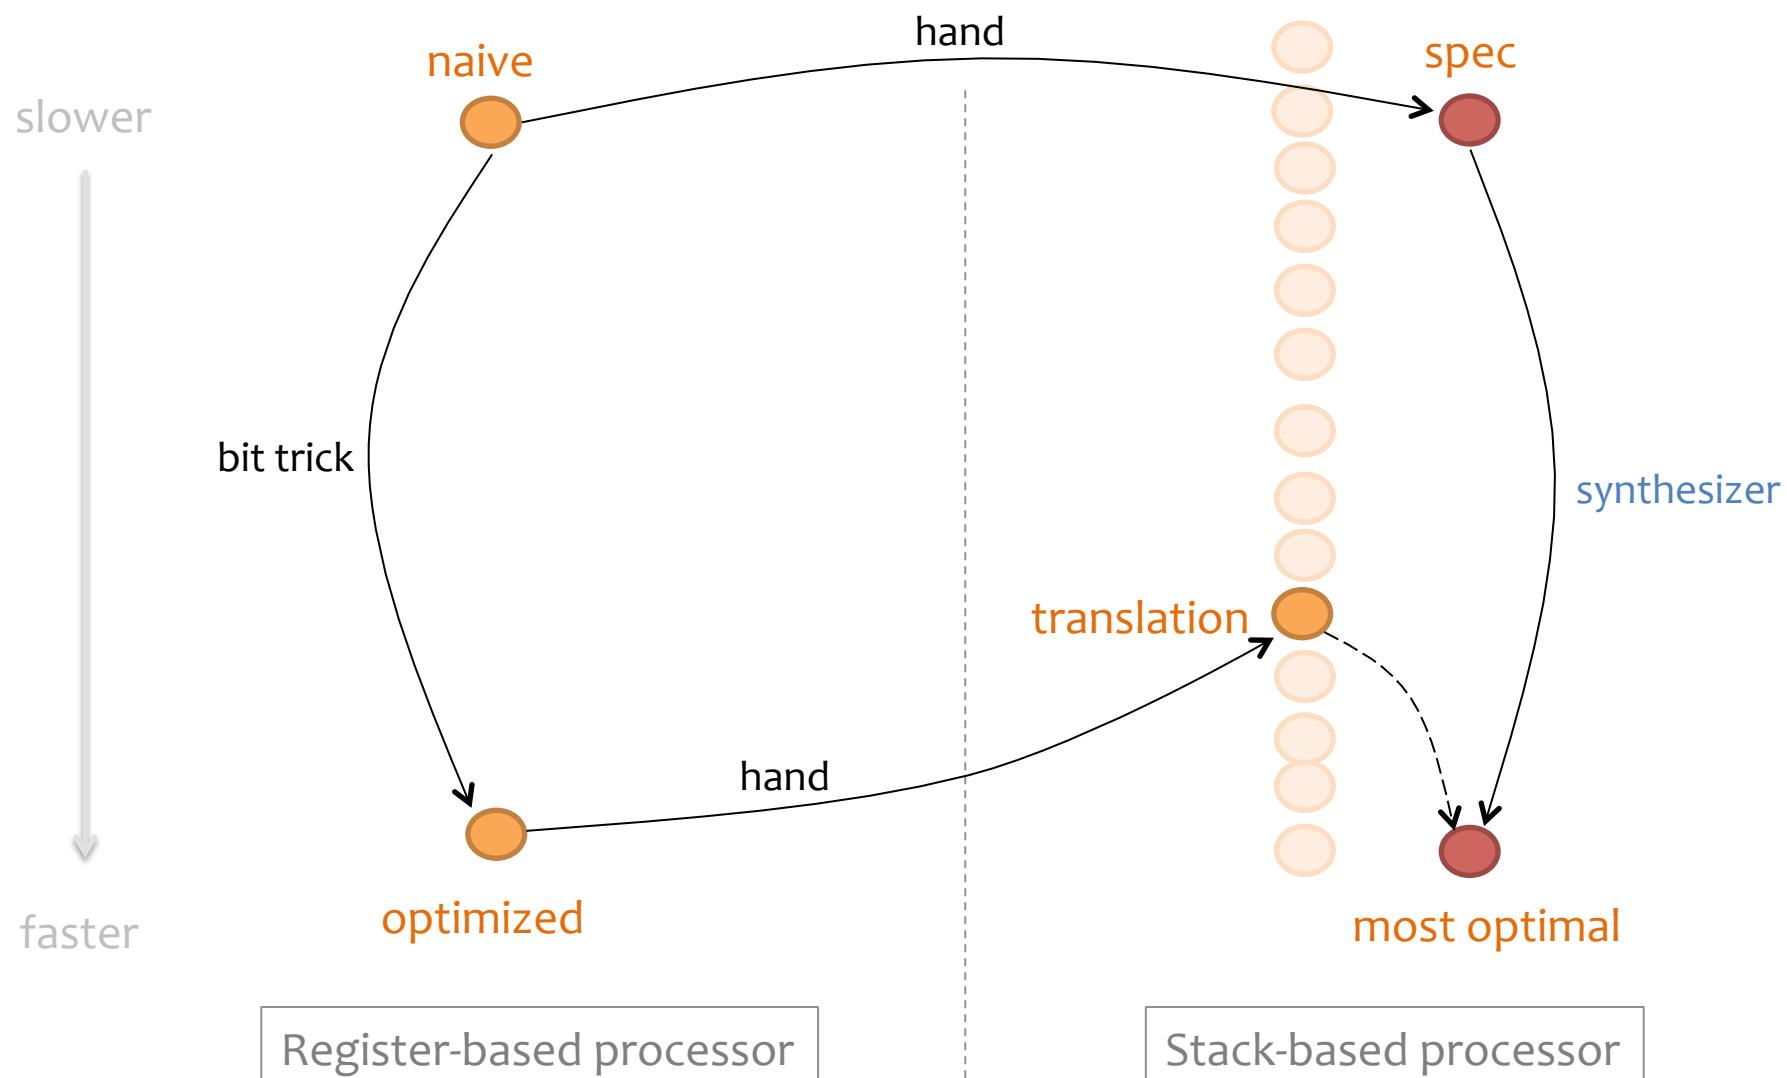
Our Experiment



Our Experiment



Comparison



Preliminary Synthesis Times

Synthesizing a program with
8 unknown instructions
takes 5 second to 5 minutes

Synthesizing a program up to
~25 unknown instructions
within 50 minutes

Preliminary Results

Program	Description	Approx. Speedup	Code length reduction
$x - (x \& y)$	Exclude common bits	5.2x	4x
$\sim(x - y)$	Negate difference	2.3x	2x
$x y$	Inclusive or	1.8x	1.8x
$(x + 7) \& -8$	Round up to multiple of 8	1.7x	1.8x
$(x \& m) (y \& \sim m)$	Replace x with y where bits of m are 1's	2x	2x
$(y \& m) (x \& \sim m)$	Replace y with x where bits of m are 1's	2.6x	2.6x
$x' = (x \& m) (y \& \sim m)$ $y' = (y \& m) (x \& \sim m)$	Swap x and y where bits of m are 1's	2x	2x

Code Length

Program	Original Length	Output Length
$x - (x \& y)$	8	2
$\sim(x - y)$	8	4
$x y$	27	15
$(x + 7) \& -8$	9	5
$(x \& m) (y \& \sim m)$	22	11
$(y \& m) (x \& \sim m)$	21	8
$x' = (x \& m) (y \& \sim m)$ $y' = (y \& m) (x \& \sim m)$	43	21

2) Synthesizing optimal library code

Input:

Sketch: program with holes to be filled

Spec: program in any programming language

Output:

Complete program with filled holes

Example: Integer Division by Constant

Naïve Implementation:

Subtract divisor until remainder < divisor.

of iterations = output value **Inefficient!**

Better Implementation:

$$\text{quotient} = (M * n) \gg s$$

- n - input
- M - “magic” number
- s - shifting value

M and s depend on the number of bits and constant divisor.

Example: Integer Division by 3

Spec:

$$n/3$$

Sketch:

$$\text{quotient} = (\text{??} * n) \gg \text{??}$$

Preliminary Results

Program	Solution	Synthesis Time (s)	Verification Time (s)	# of Pairs
$x/3$	$(43691 * x) >> 17$	2.3	7.6	4
$x/5$	$(209716 * x) >> 20$	3	8.6	6
$x/6$	$(43691 * x) >> 18$	3.3	6.6	6
$x/7$	$(149797 * x) >> 20$	2	5.5	3
deBruijn: $\log_2 x$ (x is power of 2)	deBruijn = 46, Table = $\{7, 0, 1, 3, 6, 2, 5, 4\}$	3.8	N/A	8

Note: these programs work for 18-bit number except Log2x is for 8-bit number.

3) Communication Code for GreenArray

Synthesize communication code between nodes

Interleave communication code with computational code such that

There is no deadlock.

The runtime or energy consumption of the synthesized program is minimized.

The key to the success of
low-power computing lies in
inventing better and more effective
programming frameworks.